
TimeXL: Explainable Multi-modal Time Series Prediction with LLM-in-the-Loop

Yushan Jiang^{1*}, Wenchao Yu^{2†}, Geon Lee³, Dongjin Song^{1†}, Kijung Shin³, Wei Cheng², Yanchi Liu², Haifeng Chen²

¹School of Computing, University of Connecticut

²Data Science & System Security Department, NEC Labs America

³Kim Jaechul Graduate School of AI, KAIST

Abstract

Time series analysis provides essential insights for real-world system dynamics and informs downstream decision-making, yet most existing methods often overlook the rich contextual signals present in auxiliary modalities. To bridge this gap, we introduce TimeXL, a multi-modal prediction framework that integrates a prototype-based time series encoder with three collaborating Large Language Models (LLMs) to deliver more accurate predictions and interpretable explanations. First, a multi-modal prototype-based encoder processes both time series and textual inputs to generate preliminary forecasts alongside case-based rationales. These outputs then feed into a prediction LLM, which refines the forecasts by reasoning over the encoder’s predictions and explanations. Next, a reflection LLM compares the predicted values against the ground truth, identifying textual inconsistencies or noise. Guided by this feedback, a refinement LLM iteratively enhances text quality and triggers encoder retraining. This closed-loop workflow—prediction, critique (reflect), and refinement—continuously boosts the framework’s performance and interpretability. Empirical evaluations on four real-world datasets demonstrate that TimeXL achieves up to 8.9% improvement in AUC and produces human-centric, multi-modal explanations, highlighting the power of LLM-driven reasoning for time series prediction.

1 Introduction

In the modern big-data era, time series analysis has become indispensable for understanding real-world system behaviors and guiding downstream decision-making tasks across numerous domains, including healthcare, traffic, finance, and weather [1, 2, 3, 4]. Although deep learning models have demonstrated success in capturing complex temporal dependencies [5, 6, 7, 8], real-world time series are frequently influenced by external information beyond purely temporal factors. Such additional context, which may come from textual narratives (*e.g.*, finance news [9] or medical reports [10]), can offer critical insights for more accurate forecasting and explainability.

Recent multi-modal approaches for time series have shown promise by integrating rich contextual signals from disparate data sources, thereby improving performance across a range of tasks including forecasting, classification, imputation, and retrieval [11, 12, 13, 14]. While these approaches utilize supplementary data to enhance predictive accuracy, they often lack explicit mechanisms to systematically reason and explain about *why* or *how* contextual signals affect outcomes. This gap in interpretability poses significant barriers for high-stakes applications such as finance and healthcare, where trust and transparency are paramount.

*Most of this work was done during the internship at NEC Labs America.

†Correspondence to: Wenchao Yu <wyu@nec-labs.com>, Dongjin Song <dongjin.song@uconn.edu>.

Meanwhile, Large Language Models (LLMs) [15, 16] have risen to prominence for their remarkable ability to process and reason over textual data across domains, enabling tasks like sentiment analysis, question answering, and content generation in zero- and few-shot settings [17, 18, 19]. Recent advancements in agentic designs further augment LLM capabilities, facilitating structured reasoning and iterative decision-making for context-rich scenarios [20, 21]. Their encoded domain knowledge makes them natural candidates for supporting real-world multi-modal time series analyses, where textual context (*e.g.*, news or expert notes) plays a vital role [22, 23, 24].

Motivated by these observations, we introduce TimeXL, a novel framework that adopts a closed-loop workflow of prediction, critique (reflect), and refinement, and unifies a prototype-driven time series encoder with LLM-based reasoning to deliver both accurate and interpretable multi-modal forecasting (Figure 1). Our approach first employs a *multi-modal prototype-based encoder* to generate preliminary time series predictions alongside human-readable explanations, leveraging case-based reasoning [25, 26, 27] from both the temporal and textual modalities. These explanations not only justify the encoder’s predictions but also serve as auxiliary signals to guide an LLM-powered component that further refines the forecasts and contextual rationales.

Unlike conventional methods that merely fuse multi-modal inputs for better accuracy, TimeXL iterates between predictive and refinement phases to mitigate textual noise, fill knowledge gaps, and produce more faithful explanations. Specifically, a *reflection LLM* diagnoses potential weaknesses by comparing predictions with ground-truth signals, while a *refinement LLM* incorporates these insights to update textual inputs and prototypes iteratively. This feedback loop progressively improves both the predictive and explanatory capabilities of the entire system. Our contributions are summarized as follows:

- We present a prototype-based encoder that combines time series data with textual context, producing transparent, case-based rationales.
- We exploit the interpretative prowess of LLMs to reason over the encoder’s outputs and iteratively refine both predictions and text, leading to improved prediction accuracy and explanations.
- Experiments on four real-world benchmarks show that TimeXL consistently outperforms baselines, achieving up to a 8.9% improvement in AUC while providing faithful, human-centric multi-modal explanations. Overall, TimeXL opens new avenues for explainable multi-modal time series analysis by coupling prototype-based inference with LLM-driven reasoning.

2 Related Work

Our work is primarily related to three key areas of time series research: multi-modality, explanation, and the use of LLMs. A more detailed discussion of related methods is provided in Appendix H.

Multi-modal Time Series Analysis. In recent years, multi-modal time series analysis has gained increasing attention in diverse domains such as finance, healthcare and environmental sciences [28, 11, 13]. Multiple approaches have been proposed to model interactions across different modalities, including fusion and alignment through concatenation [29], attention mechanism [12, 30, 31], gating functions [32], and contrastive learning [33, 34]. These methods facilitate tasks beyond forecasting

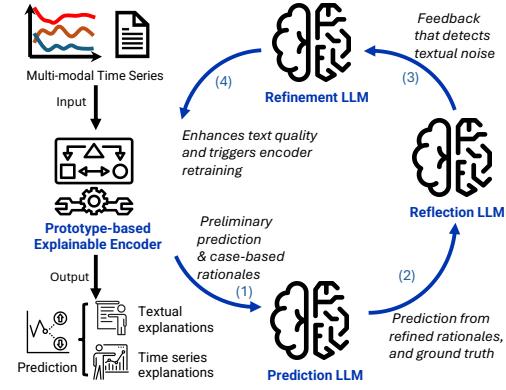


Figure 1: An overview of the TimeXL workflow. A prototype-based explainable encoder first produces predictions and case-based rationales for both time series and text. The prediction LLM refines forecasts based on these rationales (Step 1). A reflection LLM then critiques the output against ground truth (Step 2), providing feedback to detect textual noise (Step 3). Finally, a refinement LLM updates the text accordingly, triggering encoder retraining for improved accuracy and explanations (Step 4). Note that the predictions from the encoder and the LLM are also fused to enhance overall accuracy.

and classification, such as retrieval [14], imputation [13], and causal discovery [33]. Recent research efforts have also advanced the field through comprehensive benchmarks and studies [35, 36, 37], highlighting the incorporation of new modalities to boost task performance. Meanwhile, others explore transforming time series to other modalities [38, 39, 40, 41, 42] or generating synthetic ones [32, 43] for effective analysis. Nevertheless, these studies primarily focus on improving numerical performance while the deeper and tractable rationale behind *why* or *how* the textual or other contextual signals influence time series outcomes remains underexplored.

Time Series Explanation. Recent studies have explored diverse paradigms for time series interpretability. Gradient-based and perturbation-based saliency methods, for example, highlight important features at different time steps [44, 45], where some works explicitly incorporate temporal structures into models and objectives [46, 47]. Surrogate approaches also offer global or local explanations, such as applying Shapley values to time series [48], enforcing model consistency via self-supervised objectives [49], or using information-theoretic strategies for coherent explanations [50]. In contrast to saliency or surrogate-based explanations, we adopt a *case-based reasoning* paradigm [25, 26, 27], which end-to-end generates predictions and built-in explanations from learned prototypes. Our work extends this approach to multi-modal time series by producing human-readable reasoning artifacts for both the temporal and contextual modalities.

LLMs for Time Series Analysis. The rapid development of LLMs [15, 16] has begun to inspire new directions in time series research [51, 52]. Many existing methods fine-tune pre-trained LLMs on time series tasks and achieves promising results [53, 54, 55], where textual data (*e.g.*, domain instructions, metadata, summaries) are often encoded as prefix embeddings to enrich time series representations [56, 57, 58, 59]. These techniques also contribute to the emergence of time series foundation models [55, 60, 61, 62]. An alternative line of research leverages the zero-shot or few-shot capabilities of LLMs by directly prompting pre-trained language models with text-converted time series [63] or context-laden prompts representing domain knowledge [24], often yielding surprisingly strong performance in real-world scenarios. Furthermore, LLMs can act as knowledge inference modules, synthesizing high-level patterns or explanations that augment standard time series pipelines [64, 65, 66]. Building on this trend, recent works have explored time series reasoning with LLMs by framing tasks as natural-language problems, including time series understanding & question answering [67, 68, 69, 70, 71, 72], and language-guided inference [67, 69]. Compared with existing works, our work provides a unique synergy between an explainable model and LLMs, enhancing explainable prediction through iterative, grounded, and reflective interaction.

3 Methodology

In this section, we present the framework for explainable multi-modal time series prediction with LLMs. We first introduce the problem statement. Next, we present the design of a time series encoder that provides prediction and multi-modal explanations as the basis. Finally, we introduce three language agents interacting with the encoder towards better prediction and reasoning results.

3.1 Problem Statement

In this paper, we consider a multi-modal time series prediction problem. Each instance is represented by the multi-modal input (\mathbf{x}, \mathbf{s}) , where $\mathbf{x} = (x_1, x_2, \dots, x_T) \in \mathbb{R}^{N \times T}$ denotes time series data with N variables and T historical time steps, and \mathbf{s} denotes the corresponding text data describing the real-world context. The text data \mathbf{s} can be further divided into L meaningful segments. Based on the historical time series and textual context, our objective is to predict the future outcome \mathbf{y} , either as a discrete value for classification tasks, or as a continuous value for regression tasks. We focus on the classification task to reflect discrete decision-making scenarios common in real-world multi-modal applications, where interpretability is often essential for reliable decision support. We also include an evaluation of the regression setting in the Appendix G. There are four major components in the proposed TimeXL framework, a multi-modal prototype encoder \mathcal{M}_{enc} that provides an initial prediction and case-based explanations, a prediction LLM $\mathcal{M}_{\text{pred}}$ that provides a prediction based on contextual understanding with explanations, a reflection LLM $\mathcal{M}_{\text{refl}}$ that generates feedback, and a refinement LLM $\mathcal{M}_{\text{refine}}$ that refines the textual context based on the feedback. Below, we introduce each component and how they synergize toward better prediction and explanation.

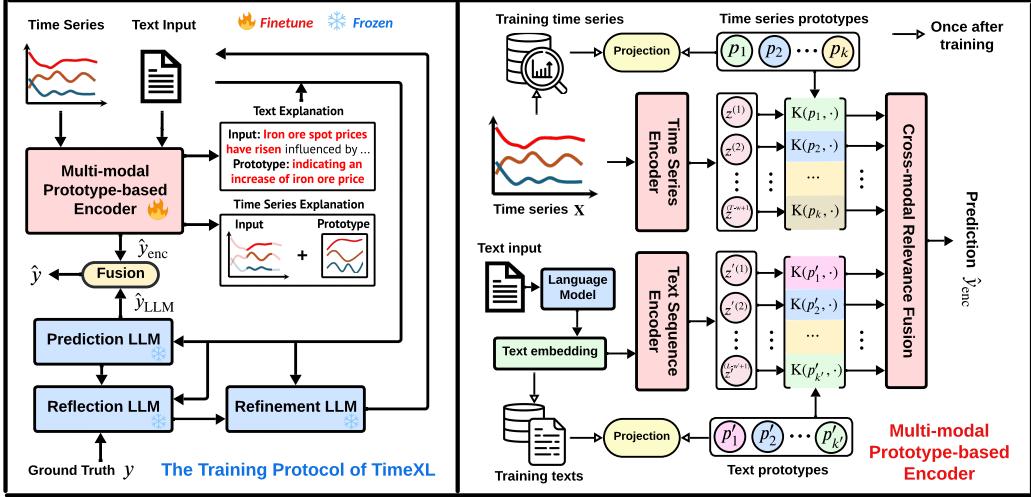


Figure 2: The training protocol of TimeXL (left), and multi-modal prototype-based encoder (right).

3.2 Multi-modal Prototype-based Encoder

We design a multi-modal prototype-based encoder that can generate predictions and explanations across different modalities in an end-to-end manner, as shown in Figure 2. We introduce the model architecture, the learning objectives that yield good explanation properties of prototypes, and the pipeline of case-based explanations using prototypes.

3.2.1 Multi-modal Sequence Modeling with Prototypes

Sequence Encoder. To capture both temporal and semantic dependencies, we adopt separate encoders for time series (\mathcal{E}_θ) and text (\mathcal{E}_ϕ). For $x \in \mathbb{R}^{N \times T}$, the time series encoder \mathcal{E}_θ maps the entire sequence into one or multiple representations, which serve as candidates for prototype learning. Simultaneously, the text input s is first transformed by a *frozen* pre-trained language model, PLM (e.g., BERT [73] or Sentence-BERT [74]), to produce embeddings $e_s \in \mathbb{R}^{d_s \times L}$. These embeddings are then processed by a separate encoder \mathcal{E}_ϕ to extract meaningful text features. It is worth noting that the choice of \mathcal{E}_θ and \mathcal{E}_ϕ also affects the granularity of explanations. As we will introduce shortly, the prototypes are learned based on sequence representations and are associated with the counterparts in the input space, where the correspondences are determined by the encoders. In this paper, we choose convolution-based encoders for both modalities to capture the fine-grained sub-sequence (*i.e.*, segment) patterns:

$$Z_{\text{time}} = (z_1, \dots, z_{T-w+1}) = \mathcal{E}_\theta(x), \quad Z_{\text{text}} = (z'_1, \dots, z'_{L-w'+1}) = \mathcal{E}_\phi(e_s) \quad (1)$$

where $z_i \in \mathbb{R}^h$ and $z'_j \in \mathbb{R}^{h'}$ denote segment-level representations learned via convolutional kernels of sizes w and w' , respectively.

Prototype Allocation. To establish interpretability, we learn a set of *time series prototypes* and *text prototypes* for each class $c \in \{1, \dots, C\}$. Specifically, we introduce: $P_{\text{time}}^{(c)} \in \mathbb{R}^{k \times h}, P_{\text{text}}^{(c)} \in \mathbb{R}^{k' \times h'}$ so that each prototype $p_i^{(c)} \in \mathbb{R}^h$ (time series) or $p'_i^{(c)} \in \mathbb{R}^{h'}$ (text) resides in the same feature space as the relevant encoder outputs. For an input sequence, we measure the similarity between each prototype and the most relevant segment in the corresponding modality:

$$\text{Sim}_i^{(c)} = \max(\text{Sim}_{i,1}^{(c)}, \dots, \text{Sim}_{i,T-w+1}^{(c)}), \text{ where } \text{Sim}_{i,j}^{(c)} = \exp\left(-\left\|p_i^{(c)} - z_j\right\|_2^2\right) \in [0, 1] \quad (2)$$

We aggregate similarity scores across all prototypes for each modality, yielding $\text{Sim}_{\text{time}} \in \mathbb{R}^{kC}$ and $\text{Sim}_{\text{text}} \in \mathbb{R}^{k'C}$. Finally, we jointly consider the cross-modal relevance and use a non-negative fusion weight matrix $W \in \mathbb{R}^{C \times (k+k')}$ that translates these scores into class probabilities:

$$\hat{y}_{\text{enc}} = \text{Softmax}\left(W[\text{Sim}_{\text{time}} \parallel \text{Sim}_{\text{text}}]\right) \in [0, 1]^C. \quad (3)$$

3.2.2 Learning Prototypes toward Better Explanation

Learning Objectives. The learning objectives include three regularization terms that reinforce the interpretability of multi-modal prototypes. In this paper, we focus on a predicting discrete label, where the basic objective is the cross-entropy loss for the prediction drawn from multi-modal explainable artifacts $\mathcal{L}_{CE} = \sum_{x,s,y} y \log(\hat{y}_{enc}) + (1-y) \log(1 - \hat{y}_{enc})$. Besides, we encourage a clustering structure of segments in the representation space by enforcing each segment representation to be adjacent to its closest prototype. Reversely, we regularize each prototype to be as close to a segment representation as possible, to help the prototype locate the most evidencing segment. Both regularization terms are denoted as \mathcal{L}_c and \mathcal{L}_e , respectively, where we omit the modality and class notations for ease of understanding:

$$\mathcal{L}_c = \sum_{z_j \in Z_{(\cdot)}} \min_{p_i \in P_{(\cdot)}} \|z_j - p_i\|_2^2, \quad \mathcal{L}_e = \sum_{p_i \in P_{(\cdot)}} \min_{z_j \in Z_{(\cdot)}} \|p_i - z_j\|_2^2 \quad (4)$$

Moreover, we encourage a diverse structure of prototype representations to avoid redundancy and maintain a compact explanation space, by penalizing their similarities via a hinge loss \mathcal{L}_d , with a threshold d_{min} :

$$\mathcal{L}_d = \sum_{i=1} \sum_{j \neq i} \max \left(0, d_{min} - \|p_i - p_j\|_2^2 \right) \quad (5)$$

The full objective is written as: $\mathcal{L} = \mathcal{L}_{CE} + \lambda_1 \mathcal{L}_c + \lambda_2 \mathcal{L}_e + \lambda_3 \mathcal{L}_d$, with hyperparameters λ_1, λ_2 , and λ_3 that balance regularization terms towards achieving an optimal and explainable prediction.

Prototype Projection. After learning objectives converge, the multi-modal prototypes are well-regularized and reflect good explanation properties. However, these prototypes are still not readily explainable as they are only close to some exemplar segments in the representation space. Therefore, we perform prototype projection to associate each prototype with a training segment from its own class that preserves \mathcal{L}_e in the representation space, for both time series and text:

$$p_i^{(c)} \leftarrow \arg \min_{z_j \in Z_{(\cdot)}^{(c)}} \left\| p_i^{(c)} - z_j \right\|_2^2, \quad \forall p_i^{(c)} \in P_{(\cdot)}^{(c)} \quad (6)$$

By associating each prototype with a training segment in the representation space, the multi-modal physical meaning is induced. During testing phase, a multi-modal instance will be compared with prototypes across different modalities to infer predictions, where the similarity scores, contribution weights, and prototypes' class information assemble the explanation artifacts for reasoning.

3.3 Explainable Prediction with LLM-in-the-Loop

To further leverage the reasoning and inference capabilities of LLMs in real-world time series contexts, we propose a framework with three interacting LLM agents: a prediction agent \mathcal{M}_{pred} , a reflection agent \mathcal{M}_{refl} , and a refinement agent \mathcal{M}_{refine} . These LLM agents interact with the multi-modal prototype-based encoder \mathcal{M}_{enc} toward better prediction accuracy and explainability.

3.3.1 Model Synergy for Augmented Prediction

Prediction with Enriched Contexts. The prediction LLM agent \mathcal{M}_{pred} generates predictions based on the input text s . To improve prediction accuracy, the encoder \mathcal{M}_{enc} supplements s with *case-based explanations*. Specifically, \mathcal{M}_{enc} selects the ω prototypes that exhibit the highest relevance to any of the textual segments within s . Relevance is determined by the similarity scores used in Equation 2. These selected prototypes are then added to the input prompt of \mathcal{M}_{pred} as explanations, providing richer real-world context and leading to more accurate predictions. The ω prototype-segment pairs, which construct the explanation expl_s of the input text s , are retrieved as follows:

$$\text{expl}_s = \left\{ \left(p_i^{(c)}, s_j \right) : (i, j, c) \in \text{Top-}\omega(\text{Sim}_{text}) \right\}, \quad \text{Top-}\omega(\text{Sim}_{text}) = \arg \text{Top-}\omega_{(i,j,c)} \left(\text{Sim}'_{i,j}^{(c)} \right)$$

Note that i, j, c denotes the prototype index, segment index, and class index, respectively. As expl_s can contain relevant contextual guidance across multiple classes, it augments the input space and removes semantic ambiguity for prediction agent \mathcal{M}_{pred} . Therefore, the prediction is drawn as $\hat{y}_{LLM} = \mathcal{M}_{pred}(s, \text{expl}_s)$. The prompt for querying \mathcal{M}_{pred} is provided in Appendix E, Figure 14.

Fused Predictions. We compile the final prediction based on a fusion of both the multi-modal encoder \mathcal{M}_{enc} and prediction LLM $\mathcal{M}_{\text{pred}}$. Specifically, we linearly combine the continuous prediction probabilities $\hat{\mathbf{y}}_{\text{enc}}$ and discrete prediction $\hat{\mathbf{y}}_{\text{LLM}}$: $\hat{\mathbf{y}} = \alpha \hat{\mathbf{y}}_{\text{enc}} + (1 - \alpha) \hat{\mathbf{y}}_{\text{LLM}}$, where $\alpha \in [0, 1]$ is the hyperparameter selected from validation data. The encoder \mathcal{M}_{enc} and prediction agent $\mathcal{M}_{\text{pred}}$ enhance each other based on their unique strengths. The \mathcal{M}_{enc} is fine-tuned based on explicit supervised signals, ensuring accuracy in capturing temporal and contextual dependencies of multi-modal time series. On the other hand, $\mathcal{M}_{\text{pred}}$ contributes deep semantic understanding drawn from extensive text corpora. By fusing predictions from two distinct perspectives, we achieve a synergistic augmentation toward more accurate predictions for complex multi-modal time series.

3.3.2 Iterative Context Refinement via Reflective Feedback

While the prediction agent $\mathcal{M}_{\text{pred}}$ leverages the explainable artifacts to make informed predictions, it is not inherently designed to fit into the context of multi-modal time series data, which could lead to inaccurate predictions when the quality of textual context is inferior. To tackle this issue, we exploit another two language agents $\mathcal{M}_{\text{refl}}$ and $\mathcal{M}_{\text{refine}}$ to generate reflective feedback and refinements on the context, respectively, for better predictive insights.

Given the prediction $\hat{\mathbf{y}}_{\text{LLM}}$ generated by the prediction agent $\mathcal{M}_{\text{pred}}$, the reflection agent $\mathcal{M}_{\text{refl}}$ aims to understand the reasoning behind the implicit prediction logic of $\mathcal{M}_{\text{pred}}$. Specifically, it generates a *reflective feedback*, Refl , by analyzing the input text s and its prediction $\hat{\mathbf{y}}_{\text{LLM}}$, against the ground truth y , to provide actionable insights for refinement, *i.e.*, $\text{Refl} = \mathcal{M}_{\text{refl}}(y, \hat{\mathbf{y}}_{\text{LLM}}, s)$. Guided by the feedback, the refinement agent $\mathcal{M}_{\text{refine}}$ refines the previous text s_i into s_{i+1} by selecting and emphasizing the most relevant content, ensuring that important patterns are appropriately contextualized, which is similar to how a domain expert would perform, *i.e.*, $s_{i+1} = \mathcal{M}_{\text{refine}}(\text{Refl}, s_i)$. The prompts for querying $\mathcal{M}_{\text{refl}}$ and $\mathcal{M}_{\text{refine}}$ are provided in Figures 16, 17 18, 19, and discussed in Appendix E.

We finally integrate the refinement via reflection into the optimization loop of our proposed TimeXL, which is summarized in Algorithm 1. Once the textual context is improved, it is used to retrain the multi-modal prototype-based encoder \mathcal{M}_{enc} for the next iteration. As such, the explanation (*e.g.*, quality of the prototypes) and predictive performance of \mathcal{M}_{enc} can be improved through this iterative process. Consequently, the prediction agent $\mathcal{M}_{\text{pred}}$ could yield better prediction with more informative inputs, further enhancing the accuracy of $\hat{\mathbf{y}}$. We evaluate the predictive performance of LLM and save the encoder and reflection when an improvement is observed ($\text{Eval}(\cdot)$ pass) when applied to the validation set. In the testing phase, we use the reflection Refl generated in the iteration with the best validation performance to guide $\mathcal{M}_{\text{refine}}$ for context refinement, mimicking how an optimized deep model is applied to testing data.

Algorithm 1 Iterative Optimization Loop of TimeXL

Inputs: Multi-modal time series (\mathbf{x}, s, y) with $\mathcal{D}_{\text{train}}$, \mathcal{D}_{val} , $\mathcal{D}_{\text{test}}$, prototype-based encoder \mathcal{M}_{enc} , prediction agent $\mathcal{M}_{\text{pred}}$, reflection agent $\mathcal{M}_{\text{refl}}$, refinement agent $\mathcal{M}_{\text{refine}}$, fusion parameter α , max iteration τ , improvement check $\text{Eval}(\cdot)$

Training & Validation:

Initialize $s_0 = s$, $i = 0$, $\mathcal{M}_{\text{enc}}^* \leftarrow \emptyset$, $\text{Refl}^* \leftarrow \emptyset$

while $i < \tau$ **do**

```

Train  $\mathcal{M}_{\text{enc}}$  using  $\mathcal{D}_{\text{train}, i} = \{(\mathbf{x}, s_i, y), \dots\}$ 
 $\hat{\mathbf{y}}_{\text{enc}}, \text{expl}_{s_i} = \mathcal{M}_{\text{enc}}(\mathbf{x}, s_i)$ 
 $\hat{\mathbf{y}}_{\text{LLM}} = \mathcal{M}_{\text{pred}}(s_i, \text{expl}_{s_i})$ 
 $\hat{\mathbf{y}} = \alpha \hat{\mathbf{y}}_{\text{enc}} + (1 - \alpha) \hat{\mathbf{y}}_{\text{LLM}}$ 
 $\text{Refl} = \mathcal{M}_{\text{refl}}(y, \hat{\mathbf{y}}_{\text{LLM}}, s_i)$ 
 $s_{i+1} = \mathcal{M}_{\text{refine}}(\text{Refl}, s_i)$ 
if  $\text{Eval}(\text{refl}, \mathcal{D}_{\text{val}})$  pass then
     $\mathcal{M}_{\text{enc}}^* \leftarrow \mathcal{M}_{\text{enc}}$ ,  $\text{Refl}^* \leftarrow \text{Refl}$ 
    increment  $i$ 

```

return $\mathcal{M}_{\text{enc}}^*$, Refl^*

Testing:

```

 $s' = \mathcal{M}_{\text{refine}}(\text{Refl}^*, s)$  for  $\mathcal{D}_{\text{test}}$ 
 $\hat{\mathbf{y}}_{\text{enc}}, \text{expl}_{s'} = \mathcal{M}_{\text{enc}}^*(\mathbf{x}, s')$ 
 $\hat{\mathbf{y}}_{\text{LLM}} = \mathcal{M}_{\text{pred}}(s', \text{expl}_{s'})$ 
 $\hat{\mathbf{y}} = \alpha \hat{\mathbf{y}}_{\text{enc}} + (1 - \alpha) \hat{\mathbf{y}}_{\text{LLM}}$ 

```

4 Experiments

4.1 Experimental Setup

Datasets. We evaluate methods on four multi-modal time series datasets from three different real-world domains, including weather, finance, and healthcare. The detailed data statistics are summarized in Table 3 of Appendix A.1. The **Weather** dataset contains meteorological reports and the hourly time series records of temperature, humidity, air pressure, wind speed, and wind direction in New York City. The task is to predict if it will rain in the next 24 hours, given the last 24 hours of weather

Table 1: The F1 score (F1) and AUROC (AUC) for TimeXL and state-of-the-art baselines on multi-modal time series datasets.

| Datasets → Methods ↓ | Weather | | Finance | | Healthcare (TP) | | Healthcare (MT) | |
|-----------------------------|--------------|--------------|--------------|--------------|-----------------|--------------|-----------------|--------------|
| | F1 | AUC | F1 | AUC | F1 | AUC | F1 | AUC |
| DLinear [75] | 0.540 | 0.660 | 0.255 | 0.485 | 0.393 | 0.500 | 0.419 | 0.388 |
| Autoformer [76] | 0.546 | 0.590 | 0.565 | 0.747 | 0.774 | 0.918 | 0.683 | 0.825 |
| Crossformer [77] | 0.500 | 0.594 | 0.571 | 0.775 | 0.924 | 0.984 | 0.737 | 0.913 |
| TimesNet [78] | 0.494 | 0.594 | 0.538 | 0.756 | 0.794 | 0.867 | 0.765 | 0.944 |
| iTransformer [79] | 0.541 | 0.650 | 0.600 | 0.783 | 0.861 | 0.931 | 0.791 | 0.963 |
| TSMixer [80] | 0.488 | 0.534 | 0.465 | 0.689 | 0.770 | 0.797 | 0.808 | 0.931 |
| TimeMixer [81] | 0.577 | 0.658 | 0.571 | 0.776 | 0.822 | 0.887 | 0.824 | 0.935 |
| FreTS [82] | 0.623 | 0.688 | 0.546 | 0.737 | 0.887 | 0.950 | 0.751 | 0.762 |
| PatchTST [5] | 0.592 | 0.675 | 0.604 | 0.795 | 0.841 | 0.934 | 0.695 | 0.928 |
| LLMTime [83] | 0.587 | 0.657 | 0.519 | 0.643 | 0.802 | 0.817 | 0.769 | 0.803 |
| PromptCast [63] | 0.499 | 0.365 | 0.418 | 0.607 | 0.727 | 0.768 | 0.696 | 0.871 |
| OFA [53] | 0.501 | 0.606 | 0.512 | 0.745 | 0.774 | 0.879 | 0.851 | 0.977 |
| FSCA [58] | 0.563 | 0.647 | 0.592 | 0.790 | 0.820 | 0.891 | 0.872 | 0.977 |
| Time-LLM [56] | 0.613 | 0.699 | 0.589 | 0.792 | 0.671 | 0.864 | 0.733 | 0.912 |
| TimeCMA [57] | 0.636 | 0.731 | 0.559 | 0.727 | 0.729 | 0.828 | 0.693 | 0.843 |
| MM-iTransformer [35] | 0.608 | 0.689 | 0.605 | 0.793 | 0.926 | 0.986 | 0.901 | 0.990 |
| MM-PatchTST [35] | 0.621 | 0.718 | 0.619 | 0.812 | 0.863 | 0.968 | 0.780 | 0.929 |
| TimeCAP [65] | <u>0.668</u> | <u>0.742</u> | <u>0.611</u> | <u>0.801</u> | <u>0.954</u> | <u>0.983</u> | <u>0.942</u> | <u>0.988</u> |
| TimeXL | 0.696 | 0.808 | 0.631 | 0.797 | 0.987 | 0.996 | 0.956 | 0.997 |

records and summary. The **Finance** dataset contains the daily record of the raw material prices together with 14 related indices from January 2017 to July 2024. Given the last 5 business days of stock price data and news, the task is to predict if the target price will exhibit an increasing, decreasing, or neutral trend on the next business day. The **Healthcare** datasets include **Test-Positive (TP)** and **Mortality (MT)**, both comprising weekly records related to influenza activity. TP contains the number of respiratory specimens testing positive for Influenza A and B, while MT reports deaths caused by influenza and pneumonia. For both datasets, the task is to predict whether the respective target ratio, either the test-positive rate or the mortality ratio in the upcoming week will exceed the historical average, based on records and summaries from the previous 20 weeks.

Baselines, Evaluation Metrics, and Setup We compare TimeXL with state-of-the-art baselines for time series prediction, including Autoformer [76], Dlinear [75], Crossformer [77], TimesNet [78], PatchTST [5], iTransformer [79], FreTS [82], TSMixer [80], TimeMixer [81] (classification implemented by TSlib [78]) and LLM-based methods like LLMTime [83], PromptCast [63], OFA [53], Time-LLM [56], TimeCMA [57] and FSCA [58], where LLMTime and PromptCast don’t need fine-tuning. While some methods are primarily used for time series prediction with continuous values, they can be easily adapted for discrete value prediction. We also evaluate the multi-modal time series methods. Besides the Time-LLM, TimeCMA and FSCA where input text is used for embedding re-programming and alignment, we also evaluate Multi-modal PatchTST and Multi-modal iTransformer from [35], as well as TimeCAP [65]. We evaluate the discrete prediction via F1 score and AUROC (AUC) score, due to label imbalance in real-world time series datasets. All datasets are split 6:2:2 for training, validation, and testing. All results are averaged over five experimental runs. We alternate different embedding methods for texts based on its average length (Sentence-BERT [74] for finance dataset and BERT [73] for the others). More details are provided in Appendix A.2, A.3, A.4.

4.2 Performance Evaluation

The results of predictive performance are shown in Table 1. It is notable that multi-modal methods generally outperform time series methods across all datasets. These methods include LLM methods (*e.g.*, Time-LLM [56], TimeCMA [57], FSCA [58]) that leverage text embeddings to enhance time series predictions. Moreover, the multi-modal variants (MM-iTransformer and MM-PatchTST [35]) improve the performance of state-of-the-art time series methods, suggesting the benefits of integrating

real-world contextual information. Besides, TimeCAP [65] integrates the predictions from both modalities, further improving the predictive performance. TimeXL constantly achieves the highest F1 and AUC scores, consistently surpassing both time series and multi-modal baselines by up to 8.9% of AUC (compared to TimeCAP on the Weather dataset). This underscores the advantage of TimeXL, which synergizes multi-modal encoder with language agents to enhance interpretability and thus predictive performance in multi-modal time series.

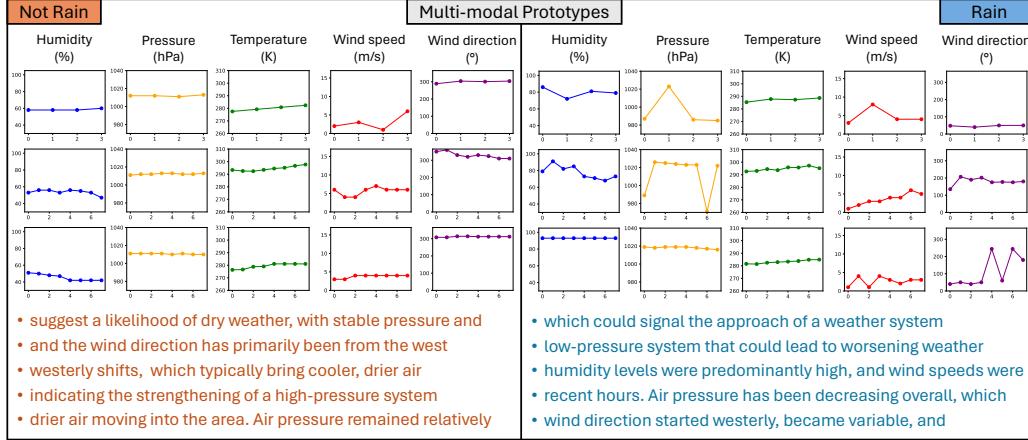


Figure 3: Key time series prototypes and text prototypes learned on the Weather dataset. Each row in the figure represents a time series prototype with different channels.

4.3 Explainable Multi-modal Prototypes

Next, we present the explainable multi-modal prototypes rendered by TimeXL, which establishes the case-based reasoning process. Figure 3 shows a subset of time series and text prototypes learned on the weather dataset. The time series prototypes demonstrate the typical temporal patterns aligned with different real-world weather conditions (*i.e.*, rain and not rain). For example, a constant or decreasing humidity at a moderate level, combined with high and steady air pressure, typically indicates a non-rainy scenario. The consistent wind direction is also a sign of mild weather conditions. On the contrary, high humidity, low and fluctuating pressure, along with variable winds typically reveal an unstable weather system ahead. In addition to time series, the text prototypes also highlight consistent semantic patterns for different weather conditions, such as the channel-specific (*e.g.*, drier air moving into the area, strengthening of high-pressure system) and overall (*e.g.*, a likelihood of dry weather) descriptions of weather activities. In Appendix D.1, we also present more multi-modal prototypes for the weather dataset in Figure 9, for the finance dataset in Figure 10, and for healthcare datasets in Figures 11 and 12. TimeXL provides coherent and informative prototypes from the exploitation of time series and its real-world contexts, which facilitates both prediction and explanation.

4.4 Multi-modal Case-based Reasoning

Building upon the multi-modal prototypes, we present a case study on the testing set of weather data, comparing the original and TimeXL’s reasoning processes to highlight its explanatory capability, as shown in Figure 4. In this case, the original text is incorrectly predicted as not rain. We have three key observations: (1) The refinement process filters the original text to emphasize weather conditions more indicative of rain, guided by reflections from training examples. The refined text preserves the statement on stability while placing more emphasis on humidity and wind as key indicators. (2) Accordingly, the matched segment-prototype pairs from the original text focus more on temperature stability and typical diurnal variations, while the matched pairs in the refined text highlights wind variability, moisture transport, and approaching weather system, aligning more with rain conditions. (3) Furthermore, the reasoning on time series provides a complementary view for assessing weather conditions. The matched time series prototypes identify high humidity and its drop-and-rise trends, wind speed fluctuations and directional shifts, and the declining phase of air pressure fluctuations,

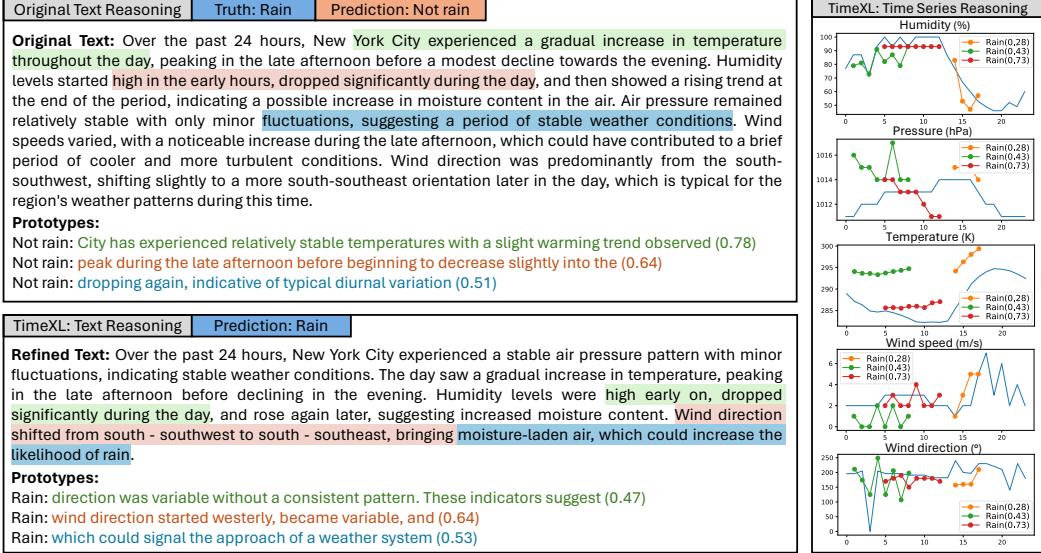


Figure 4: Multi-modal case-based reasoning example on the Weather dataset. The left part illustrates the reasoning process for both the original and refined text in TimeXL, with matched prototype-input pairs highlighted in the same color along with their similarity scores. The right part presents the time series reasoning in TimeXL, where matched prototypes are overlaid on the time series.

all of which are linked to the upcoming rainy conditions. The matched multi-modal prototypes from TimeXL demonstrate its effectiveness in capturing relevant information for both predictive and explanatory analysis. We also provide a case study on the Finance dataset in Figure 13, where textual explanations are generated at the granularity of a half-sentence.

4.5 Iterative Analysis

To verify the effectiveness of overall workflow with reflection and refinement LLMs as shown in Figure 1, we conduct an iterative analysis of text quality and TimeXL performance, as shown in Figure 5. Specifically, we evaluate the text quality based on its zero-shot predictive accuracy using an LLM. Notably, the text quality benefits from iteration improvements and mostly saturates after one or two iterations. Correspondingly, TimeXL performance quickly improves and stabilizes with very minor fluctuations. These observations underscore how TimeXL alternates between predictive and reflective refinement phases to mitigate textual noise, thus enhancing its predictive capability.

4.6 Ablation Studies

We present the component ablations of TimeXL in Table 2, where the texts are refined using the reflections with the best validation performance. Firstly, the multi-modal encoder consistently outperforms single-modality variants across all datasets, highlighting the benefit of integrating time series and textual modalities. Secondly, the prediction LLM exhibits better performance than text-only encoder, underscoring the advantage of LLM’s contextual understanding and reasoning. Furthermore, the text prototypes consistently improve the predictive performance of LLM, underscoring the effectiveness of explainable artifacts from the multi-modal encoder, in terms of providing relevant contextual guidance. Finally, the fusion of prediction LLM and multi-modal encoder further enhances

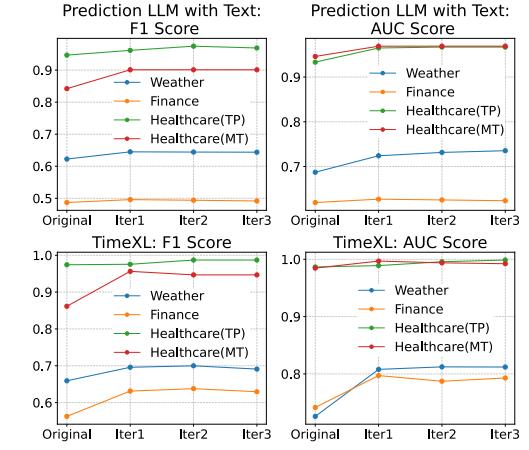


Figure 5: The text quality and TimeXL performance over iterations.

Table 2: Ablation studies of TimeXL on multi-modal time series datasets.

| Ablation ↓ | Variants | Weather | | Finance | | Healthcare (Test-Positive) | | Healthcare (Mortality) | |
|------------|------------------|--------------|--------------|--------------|--------------|----------------------------|--------------|------------------------|--------------|
| | | F1 | AUC | F1 | AUC | F1 | AUC | F1 | AUC |
| Encoder | Time Series | 0.602 | 0.691 | 0.585 | 0.751 | 0.889 | 0.957 | 0.861 | 0.965 |
| | Text | 0.567 | 0.658 | 0.472 | 0.636 | 0.871 | 0.941 | 0.840 | 0.887 |
| | Multi-modal | 0.674 | 0.767 | 0.619 | 0.791 | 0.934 | 0.974 | 0.937 | 0.988 |
| LLM | Text | 0.645 | 0.724 | 0.496 | 0.627 | 0.974 | 0.967 | 0.901 | 0.969 |
| | Text + Prototype | 0.667 | 0.739 | 0.544 | 0.662 | 0.987 | 0.983 | 0.952 | 0.976 |
| Fusion | Select-Best | 0.674 | 0.767 | 0.619 | 0.791 | 0.987 | 0.983 | 0.952 | 0.988 |
| | TimeXL | 0.696 | 0.808 | 0.631 | 0.797 | 0.987 | 0.996 | 0.956 | 0.997 |

the predictive performance, surpassing the best results achieved by either component alone. These observations demonstrate the advantage of our framework synergizing the time series model and LLM for mutually augmented prediction. In Appendix B, additional ablation and model studies are presented, including analyses of the multi-modal encoder (Figure 6; Table 4, 5, 6), quantitative evaluation of case-based explanations (Figure 7), and comparisons of base LLMs (Table 7; Figure 8).

5 Conclusions

In this paper, we present TimeXL, an explainable multi-modal time series prediction framework that synergizes a designed prototype-based encoder with three collaborative LLM agents in the loop (prediction, reflection, and refinement) to deliver more accurate predictions and explanations. Experiments on four multi-modal time series datasets show the advantages of TimeXL over state-of-the-art baselines and its excellent explanation capabilities.

6 Acknowledgments

Dongjin Song gratefully acknowledges the support of the National Science Foundation under Grant No. 2338878, as well as the generous research gifts from NEC Labs America and Morgan Stanley.

References

- [1] Bo Jin, Haoyu Yang, Leilei Sun, Chuanren Liu, Yue Qu, and Jianing Tong. A treatment engine by predicting next-period prescriptions. In *Proceedings of the 24th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, pages 1608–1616, 2018.
- [2] Shengnan Guo, Youfang Lin, Ning Feng, Chao Song, and Huaiyu Wan. Attention based spatial-temporal graph convolutional networks for traffic flow forecasting. In *Proceedings of the AAAI conference on artificial intelligence*, volume 33, pages 922–929, 2019.
- [3] Liheng Zhang, Charu Aggarwal, and Guo-Jun Qi. Stock price prediction via discovering multi-frequency trading patterns. In *Proceedings of the 23rd ACM SIGKDD international conference on knowledge discovery and data mining*, pages 2141–2149, 2017.
- [4] Yao Qin, Dongjin Song, Haifeng Cheng, Wei Cheng, Guofei Jiang, and Garrison W Cottrell. A dual-stage attention-based recurrent neural network for time series prediction. In *Proceedings of the 26th International Joint Conference on Artificial Intelligence*, pages 2627–2633, 2017.
- [5] Yuqi Nie, Nam H Nguyen, Phanwadee Sinthong, and Jayant Kalagnanam. A time series is worth 64 words: Long-term forecasting with transformers. In *The Eleventh International Conference on Learning Representations*, 2023.
- [6] Ailin Deng and Bryan Hooi. Graph neural network-based anomaly detection in multivariate time series. In *Proceedings of the AAAI conference on artificial intelligence*, volume 35, pages 4027–4035, 2021.
- [7] Xiang Zhang, Ziyuan Zhao, Theodoros Tsiligkaridis, and Marinka Zitnik. Self-supervised contrastive pre-training for time series via time-frequency consistency. *Advances in Neural Information Processing Systems*, 35:3988–4003, 2022.

- [8] Yong Liu, Chenyu Li, Jianmin Wang, and Mingsheng Long. Koopa: Learning non-stationary time series dynamics with koopman predictors. *Advances in Neural Information Processing Systems*, 36, 2024.
- [9] Zihan Dong, Xinyu Fan, and Zhiyuan Peng. Fnspid: A comprehensive financial news dataset in time series. In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 4918–4927, 2024.
- [10] Ryan King, Tianbao Yang, and Bobak J Mortazavi. Multimodal pretraining of medical time series and notes. In *Machine Learning for Health (ML4H)*, pages 244–255. PMLR, 2023.
- [11] Ke Niu, Ke Zhang, Xueping Peng, Yijie Pan, and Naian Xiao. Deep multi-modal intermediate fusion of clinical record and time series data in mortality prediction. *Frontiers in Molecular Biosciences*, 10:1136071, 2023.
- [12] Geon Lee, Wenchao Yu, Wei Cheng, and Haifeng Chen. Moat: Multi-modal augmented time series forecasting. 2024.
- [13] Xiaohu Zhao, Kebin Jia, Benjamin Letcher, Jennifer Fair, Yiqun Xie, and Xiaowei Jia. Vimts: Variational-based imputation for multi-modal time series. In *2022 IEEE International Conference on Big Data (Big Data)*, pages 349–358. IEEE, 2022.
- [14] Tom Bamford, Andrea Coletta, Elizabeth Fons, Sriram Gopalakrishnan, Svitlana Vytnenko, Tucker Balch, and Manuela Veloso. Multi-modal financial time-series retrieval through latent space projections. In *Proceedings of the Fourth ACM International Conference on AI in Finance*, pages 498–506, 2023.
- [15] Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altenschmidt, Sam Altman, Shyamal Anadkat, et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
- [16] Gemini Team, Rohan Anil, Sebastian Borgeaud, Yonghui Wu, Jean-Baptiste Alayrac, Jiahui Yu, Radu Soricut, Johan Schalkwyk, Andrew M Dai, Anja Hauth, et al. Gemini: a family of highly capable multimodal models. *arXiv preprint arXiv:2312.11805*, 2023.
- [17] Wenxuan Zhang, Yue Deng, Bing Liu, Sinno Pan, and Lidong Bing. Sentiment analysis in the era of large language models: A reality check. In *Findings of the Association for Computational Linguistics: NAACL 2024*, pages 3881–3906, 2024.
- [18] Ehsan Kamalloo, Nouha Dziri, Charles Clarke, and Davood Rafiei. Evaluating open-domain question answering in the era of large language models. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 5591–5606, 2023.
- [19] Zixuan Wang, Qinkai Duan, Yu-Wing Tai, and Chi-Keung Tang. C3llm: Conditional multimodal content generation using large language models. *arXiv preprint arXiv:2405.16136*, 2024.
- [20] Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: Language agents with verbal reinforcement learning. *Advances in Neural Information Processing Systems*, 36, 2024.
- [21] Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegreffe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, et al. Self-refine: Iterative refinement with self-feedback. *Advances in Neural Information Processing Systems*, 36, 2024.
- [22] Xin Liu, Daniel McDuff, Geza Kovacs, Isaac Galatzer-Levy, Jake Sunshine, Jiening Zhan, Ming-Zher Poh, Shun Liao, Paolo Di Achille, and Shwetak Patel. Large language models are few-shot health learners.
- [23] Kelvin JL Koa, Yunshan Ma, Ritchie Ng, and Tat-Seng Chua. Learning to generate explainable stock predictions using self-reflective large language models. In *Proceedings of the ACM on Web Conference 2024*, pages 4304–4315, 2024.

- [24] Xinglei Wang, Meng Fang, Zichao Zeng, and Tao Cheng. Where would i go next? large language models as human mobility predictors. *arXiv preprint arXiv:2308.15197*, 2023.
- [25] Yao Ming, Panpan Xu, Huamin Qu, and Liu Ren. Interpretable and steerable sequence learning via prototypes. In *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, pages 903–913, 2019.
- [26] Jingchao Ni, Zhengzhang Chen, Wei Cheng, Bo Zong, Dongjin Song, Yanchi Liu, Xuchao Zhang, and Haifeng Chen. Interpreting convolutional sequence model by learning local prototypes with adaptation regularization. In *Proceedings of the 30th ACM International Conference on Information & Knowledge Management*, pages 1366–1375, 2021.
- [27] Yushan Jiang, Wenchao Yu, Dongjin Song, Lu Wang, Wei Cheng, and Haifeng Chen. Fedskill: Privacy preserved interpretable skill learning via imitation. In *Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 1010–1019, 2023.
- [28] Geri Skenderi, Christian Joppi, Matteo Denitto, and Marco Cristani. Well googled is half done: Multimodal forecasting of new fashion product sales with image-based google trends. *Journal of Forecasting*, 43(6):1982–1997, 2024.
- [29] Zihao Li, Xiao Lin, Zhining Liu, Jiaru Zou, Ziwei Wu, Lecheng Zheng, Dongqi Fu, Yada Zhu, Hendrik Hamann, Hanghang Tong, and Jingrui He. Language in the flow of time: Time-series-paired texts weaved into a unified temporal narrative, 2025.
- [30] Sameep Chattpadhyay, Pulkit Paliwal, Sai Shankar Narasimhan, Shubhankar Agarwal, and Sandeep P Chinchali. Context matters: Leveraging contextual features for time series forecasting. *arXiv preprint arXiv:2410.12672*, 2024.
- [31] Chen Su, Yuanhe Tian, and Yan Song. Multimodal conditioned diffusive time series forecasting. *arXiv preprint arXiv:2504.19669*, 2025.
- [32] Siru Zhong, Weilin Ruan, Ming Jin, Huan Li, Qingsong Wen, and Yuxuan Liang. Time-vlm: Exploring multimodal vision-language models for augmented time series forecasting. *arXiv preprint arXiv:2502.04395*, 2025.
- [33] Lecheng Zheng, Zhengzhang Chen, Jingrui He, and Haifeng Chen. Mulan: Multi-modal causal structure learning and root cause analysis for microservice systems. In *Proceedings of the ACM on Web Conference 2024*, pages 4107–4116, 2024.
- [34] Jun Li, Che Liu, Sibo Cheng, Rossella Arcucci, and Shenda Hong. Frozen language model helps ecg zero-shot learning. In *Medical Imaging with Deep Learning*, pages 402–415. PMLR, 2024.
- [35] Haoxin Liu, Shangqing Xu, Zhiyuan Zhao, Lingkai Kong, Harshavardhan Kamarthi, Aditya B Sasanur, Megha Sharma, Jiaming Cui, Qingsong Wen, Chao Zhang, et al. Time-mmd: A new multi-domain multimodal dataset for time series analysis. *arXiv preprint arXiv:2406.08627*, 2024.
- [36] Andrew Robert Williams, Arjun Ashok, Étienne Marcotte, Valentina Zantedeschi, Jithendaraa Subramanian, Roland Riachi, James Requeima, Alexandre Lacoste, Irina Rish, Nicolas Chapados, et al. Context is key: A benchmark for forecasting with essential textual information. *arXiv preprint arXiv:2410.18959*, 2024.
- [37] Yushan Jiang, Kanghui Ning, Zijie Pan, Xuyang Shen, Jingchao Ni, Wenchao Yu, Anderson Schneider, Haifeng Chen, Yuriy Nevmyvaka, and Dongjin Song. Multi-modal time series analysis: A tutorial and survey. *arXiv preprint arXiv:2503.13709*, 2025.
- [38] Zekun Li, Shiyang Li, and Xifeng Yan. Time series as images: Vision transformer for irregularly sampled time series. *Advances in Neural Information Processing Systems*, 36:49187–49204, 2023.
- [39] Shi Bin Hoo, Samuel Müller, David Salinas, and Frank Hutter. The tabular foundation model tabPFN outperforms specialized time series forecasting models based on simple features. *arXiv preprint arXiv:2501.02945*, 2025.

- [40] Jingchao Ni, Ziming Zhao, ChengAo Shen, Hanghang Tong, Dongjin Song, Wei Cheng, Dongsheng Luo, and Haifeng Chen. Harnessing vision models for time series analysis: A survey. *arXiv preprint arXiv:2502.08869*, 2025.
- [41] Xiongxiao Xu, Yue Zhao, S Yu Philip, and Kai Shu. Beyond numbers: A survey of time series analysis in the era of multimodal llms. *Authorea Preprints*, 2025.
- [42] Mouxiang Chen, Lefei Shen, Zhuo Li, Xiaoyun Joy Wang, Jianling Sun, and Chenghao Liu. Visions: Visual masked autoencoders are free-lunch zero-shot time series forecasters. *arXiv preprint arXiv:2408.17253*, 2024.
- [43] Xu Liu, Taha Aksu, Juncheng Liu, Qingsong Wen, Yuxuan Liang, Caiming Xiong, Silvio Savarese, Doyen Sahoo, Junnan Li, and Chenghao Liu. Empowering time series analysis with synthetic data: A survey and outlook in the era of foundation models. *arXiv preprint arXiv:2503.11411*, 2025.
- [44] Aya Abdelsalam Ismail, Mohamed Gunady, Hector Corrada Bravo, and Soheil Feizi. Benchmarking deep learning interpretability in time series predictions. *Advances in neural information processing systems*, 33:6441–6452, 2020.
- [45] Sana Tonekaboni, Shalmali Joshi, Kieran Campbell, David K Duvenaud, and Anna Goldenberg. What went wrong and when? instance-wise feature importance for time-series black-box models. *Advances in Neural Information Processing Systems*, 33:799–809, 2020.
- [46] Kin Kwan Leung, Clayton Rooke, Jonathan Smith, Saba Zuberi, and Maksims Volkovs. Temporal dependencies in feature importance for time series prediction. In *The Eleventh International Conference on Learning Representations*.
- [47] Jonathan Crabbé and Mihaela Van Der Schaar. Explaining time series predictions with dynamic masks. In *International Conference on Machine Learning*, pages 2166–2177. PMLR, 2021.
- [48] João Bento, Pedro Saleiro, André F Cruz, Mário AT Figueiredo, and Pedro Bizarro. Timeshap: Explaining recurrent models through sequence perturbations. In *Proceedings of the 27th ACM SIGKDD conference on knowledge discovery & data mining*, pages 2565–2573, 2021.
- [49] Owen Queen, Tom Hartvigsen, Teddy Koker, Huan He, Theodoros Tsiligkaridis, and Marinka Zitnik. Encoding time-series explanations through self-supervised model behavior consistency. *Advances in Neural Information Processing Systems*, 36, 2024.
- [50] Zichuan Liu, Tianchun Wang, Jimeng Shi, Xu Zheng, Zhuomin Chen, Lei Song, Wenqian Dong, Jayantha Obeysekera, Farhad Shirani, and Dongsheng Luo. Timex++: Learning time-series explanations with information bottleneck. In *Forty-first International Conference on Machine Learning*, 2024.
- [51] Yushan Jiang, Zijie Pan, Xikun Zhang, Sahil Garg, Anderson Schneider, Yuriy Nevmyvaka, and Dongjin Song. Empowering time series analysis with large language models: A survey. In Kate Larson, editor, *Proceedings of the Thirty-Third International Joint Conference on Artificial Intelligence, IJCAI-24*, pages 8095–8103. International Joint Conferences on Artificial Intelligence Organization, 8 2024. Survey Track.
- [52] Yuxuan Liang, Haomin Wen, Yuqi Nie, Yushan Jiang, Ming Jin, Dongjin Song, Shirui Pan, and Qingsong Wen. Foundation models for time series analysis: A tutorial and survey. In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 6555–6565, 2024.
- [53] Tian Zhou, Peisong Niu, Liang Sun, Rong Jin, et al. One fits all: Power general time series analysis by pretrained lm. *Advances in neural information processing systems*, 36:43322–43355, 2023.
- [54] Yuxuan Bian, Xuan Ju, Jiangtong Li, Zhijian Xu, Dawei Cheng, and Qiang Xu. Multi-patch prediction: Adapting language models for time series representation learning. In *Forty-first International Conference on Machine Learning*.

- [55] Abdul Fatir Ansari, Lorenzo Stella, Caner Turkmen, Xiyuan Zhang, Pedro Mercado, Huibin Shen, Oleksandr Shchur, Syama Sundar Rangapuram, Sebastian Pineda Arango, Shubham Kapoor, et al. Chronos: Learning the language of time series. *arXiv preprint arXiv:2403.07815*, 2024.
- [56] Ming Jin, Shiyu Wang, Lintao Ma, Zhixuan Chu, James Y Zhang, Xiaoming Shi, Pin-Yu Chen, Yuxuan Liang, Yuan-Fang Li, Shirui Pan, et al. Time-llm: Time series forecasting by reprogramming large language models. In *The Twelfth International Conference on Learning Representations*, 2024.
- [57] Chenxi Liu, Qianxiong Xu, Hao Miao, Sun Yang, Lingzheng Zhang, Cheng Long, Ziyue Li, and Rui Zhao. Timecma: Towards llm-empowered multivariate time series forecasting via cross-modality alignment. In *AAAI*, 2025.
- [58] Yuxiao Hu, Qian Li, Dongxiao Zhang, Jinyue Yan, and Yuntian Chen. Context-alignment: Activating and enhancing LLMs capabilities in time series. In *The Thirteenth International Conference on Learning Representations*, 2025.
- [59] Chenxi Liu, Shaowen Zhou, Hao Miao, Qianxiong Xu, Cheng Long, Ziyue Li, and Rui Zhao. Efficient multivariate time series forecasting via calibrated language models with privileged knowledge distillation. *arXiv preprint arXiv:2505.02138*, 2025.
- [60] Abhimanyu Das, Weihao Kong, Rajat Sen, and Yichen Zhou. A decoder-only foundation model for time-series forecasting. *arXiv preprint arXiv:2310.10688*, 2023.
- [61] Gerald Woo, Chenghao Liu, Akshat Kumar, Caiming Xiong, Silvio Savarese, and Doyen Sahoo. Unified training of universal time series forecasting transformers. *arXiv preprint arXiv:2402.02592*, 2024.
- [62] Yong Liu, Haoran Zhang, Chenyu Li, Xiangdong Huang, Jianmin Wang, and Mingsheng Long. Timer: Transformers for time series analysis at scale. *arXiv preprint arXiv:2402.02368*, 2024.
- [63] Hao Xue and Flora D Salim. Promptcast: A new prompt-based learning paradigm for time series forecasting. *IEEE Transactions on Knowledge and Data Engineering*, 2023.
- [64] Zihan Chen, Lei Nico Zheng, Cheng Lu, Jialu Yuan, and Di Zhu. Chatgpt informed graph neural network for stock movement prediction. Available at SSRN 4464002, 2023.
- [65] Geon Lee, Wenchao Yu, Kijung Shin, Wei Cheng, and Haifeng Chen. Timecap: Learning to contextualize, augment, and predict time series events with large language model agents. In *AAAI*, 2025.
- [66] Xinlei Wang, Maike Feng, Jing Qiu, Jinjin Gu, and Junhua Zhao. From news to forecast: Integrating event analysis in llm-based time series forecasting with reflection. *arXiv preprint arXiv:2409.17515*, 2024.
- [67] Jialin Chen, Aosong Feng, Ziyu Zhao, Juan Garza, Gaukhar Nurbek, Cheng Qin, Ali Maatouk, Leandros Tassiulas, Yifeng Gao, and Rex Ying. Mtbench: A multimodal time series benchmark for temporal reasoning and question answering. *arXiv preprint arXiv:2503.16858*, 2025.
- [68] Yaxuan Kong, Yiyuan Yang, Shiyu Wang, Chenghao Liu, Yuxuan Liang, Ming Jin, Stefan Zohren, Dan Pei, Yan Liu, and Qingsong Wen. Position: Empowering time series reasoning with multimodal llms. *arXiv preprint arXiv:2502.01477*, 2025.
- [69] Yaxuan Kong, Yiyuan Yang, Yoontae Hwang, Wenjie Du, Stefan Zohren, Zhangyang Wang, Ming Jin, and Qingsong Wen. Time-mqa: Time series multi-task question answering with context enhancement. *arXiv preprint arXiv:2503.01875*, 2025.
- [70] Chengsen Wang, Qi Qi, Jingyu Wang, Haifeng Sun, Zirui Zhuang, Jinming Wu, Lei Zhang, and Jianxin Liao. Chattime: A unified multimodal time series foundation model bridging numerical and textual data. *arXiv preprint arXiv:2412.11376*, 2024.
- [71] Zhe Xie, Zeyan Li, Xiao He, Longlong Xu, Xidao Wen, Tieying Zhang, Jianjun Chen, Rui Shi, and Dan Pei. Chatts: Aligning time series with llms via synthetic data for enhanced understanding and reasoning. *arXiv preprint arXiv:2412.03104*, 2024.

- [72] Yifu Cai, Arjun Choudhry, Mononito Goswami, and Artur Dubrawski. Timeseriesexam: A time series understanding exam. *arXiv preprint arXiv:2410.14752*, 2024.
- [73] Jacob Devlin Ming-Wei Chang Kenton and Lee Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of naacL-HLT*, volume 1, page 2. Minneapolis, Minnesota, 2019.
- [74] Nils Reimers and Iryna Gurevych. Sentence-bert: Sentence embeddings using siamese bert-networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 11 2019.
- [75] Ailing Zeng, Muxi Chen, Lei Zhang, and Qiang Xu. Are transformers effective for time series forecasting? In *Proceedings of the AAAI conference on artificial intelligence*, volume 37, pages 11121–11128, 2023.
- [76] Haixu Wu, Jiehui Xu, Jianmin Wang, and Mingsheng Long. Autoformer: Decomposition transformers with auto-correlation for long-term series forecasting. *Advances in neural information processing systems*, 34:22419–22430, 2021.
- [77] Yunhao Zhang and Junchi Yan. Crossformer: Transformer utilizing cross-dimension dependency for multivariate time series forecasting. In *The eleventh international conference on learning representations*, 2023.
- [78] Haixu Wu, Tengge Hu, Yong Liu, Hang Zhou, Jianmin Wang, and Mingsheng Long. Timesnet: Temporal 2d-variation modeling for general time series analysis. In *The Eleventh International Conference on Learning Representations*, 2023.
- [79] Yong Liu, Tengge Hu, Haoran Zhang, Haixu Wu, Shiyu Wang, Lintao Ma, and Mingsheng Long. itransformer: Inverted transformers are effective for time series forecasting. In *The Twelfth International Conference on Learning Representations*, 2024.
- [80] Si-An Chen, Chun-Liang Li, Sercan O Arik, Nathanael Christian Yoder, and Tomas Pfister. Tsmixer: An all-mlp architecture for time series forecasting. *Transactions on Machine Learning Research*, 2023.
- [81] Shiyu Wang, Haixu Wu, Xiaoming Shi, Tengge Hu, Huakun Luo, Lintao Ma, James Y Zhang, and JUN ZHOU. Timemixer: Decomposable multiscale mixing for time series forecasting. In *The Twelfth International Conference on Learning Representations*.
- [82] Kun Yi, Qi Zhang, Wei Fan, Shoujin Wang, Pengyang Wang, Hui He, Ning An, Defu Lian, Longbing Cao, and Zhendong Niu. Frequency-domain mlps are more effective learners in time series forecasting. *Advances in Neural Information Processing Systems*, 36, 2024.
- [83] Nate Gruver, Marc Finzi, Shikai Qiu, and Andrew Gordon Wilson. Large Language Models Are Zero Shot Time Series Forecasters. In *NeurIPS*, 2023.
- [84] Diederik P Kingma and Jimmy Ba. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*, 2014.
- [85] Jiawei Gu, Xuhui Jiang, Zhichao Shi, Hexiang Tan, Xuehao Zhai, Chengjin Xu, Wei Li, Yinghan Shen, Shengjie Ma, Honghao Liu, et al. A survey on llm-as-a-judge. *arXiv preprint arXiv:2411.15594*, 2024.
- [86] Mingyue Cheng, Yiheng Chen, Qi Liu, Zhiding Liu, Yucong Luo, and Enhong Chen. Instructime: Advancing time series classification with multimodal language modeling. In *Proceedings of the Eighteenth ACM International Conference on Web Search and Data Mining*, WSDM ’25, page 792–800, New York, NY, USA, 2025. Association for Computing Machinery.

A Experimental Settings

A.1 Dataset Statistics

In this subsection, we provide more details of the real-world datasets we used for the experiments. The data statistics are summarized in Table 3, including the meta information (*e.g.*, domain resolution, duration of real-world time series records), the number of channels and timesteps and so on. We use the Weather and Healthcare datasets in TimeCAP [65], and a Finance dataset extended from [12].

The **Weather** dataset contains the hourly time series record of temperature, humidity, air pressure, wind speed, and wind direction³, and related weather summaries in New York City from October 2012 to November 2017. The task is to predict if it will rain in the next 24 hours, given the last 24 hours of weather records and summary.

The **Finance** dataset contains the daily record of the raw material prices together with 14 related indices ranging from January 2017 to July 2024⁴ with news articles. The task is to predict if future prices will increase by more than 1%, decrease by more than 1%, or exhibit a neutral trend on the next business day, given the last 5 business days of stock price data and news.

The Healthcare datasets are related to testing cases and deaths of influenza⁵. The **Healthcare (Test-Positive)** dataset consists of the weekly records of the number of positive specimens for Influenza A and B, and related healthcare reports. The task is to predict if the percentage of respiratory specimens testing positive in the upcoming week for influenza will exceed the average value, given the records and summary in the last 20 weeks. Similarly, the **Healthcare (Mortality)** dataset contains the weekly records and healthcare reports of influenza and pneumonia deaths. The task is to predict if the mortality ratio from influenza and pneumonia will exceed the average value, given the records and summary in the last 20 weeks.

Table 3: Summary of dataset statistics.

| Domain | Dataset | Resolution | # Channels | # Timesteps | Duration | Ground Truth Distribution |
|------------|---------------|------------|------------|-------------|-------------------|---|
| Weather | New York | Hourly | 5 | 45,216 | 2012.10 - 2017.11 | Rain (24.26%) / Not rain (75.74%) |
| Finance | Raw Material | Daily | 15 | 1,876 | 2012.09 - 2022.02 | Inc. (36.7%) / Dec. (34.1%) / Neutral (29.2%) |
| Healthcare | Test-Positive | Weekly | 6 | 447 | 2015.10 - 2024.04 | Not exceed (65.77%) / Exceed (34.23%) |
| Healthcare | Mortality | Weekly | 4 | 395 | 2016.07 - 2024.06 | Not exceed (69.33%) / Exceed (30.67%) |

A.2 Hyperparameters

First, we provide the hyperparameters of baseline methods. Unless otherwise specified, we used the default hyperparameters from the Time Series Library (TSLib) [78]. For LLMTTime [83], OFA [53], Time-LLM [56], TimeCMA [57], TimeCAP [65], FSCA [58], we use their own implementations. For all methods, the dropout rate $\in \{0.0, 0.1, 0.2\}$, learning rate $\in \{0.0001, 0.0003, 0.001\}$. For transformer-based and LLM fine-tuning methods [76, 77, 79, 5, 57, 56], the number of attention layers $\in \{1, 2\}$, the number of attention heads $\in \{4, 8, 16\}$. For Dlinear [75], moving average $\in \{3, 5\}$. For TimesNet [78] the number of layers $\in \{1, 2\}$. For PatchTST [5], MM-PatchTST [35] and FSCA [58], the patch size $\in \{3, 5\}$ for the finance dataset. For TimeCAP [65], we use the encoder with the prediction LLM, given the inherent multi-modal time series input. Next, we provide the hyperparameters of TimeXL. The numbers of time series prototypes and text prototypes are $k \in \{5, 10, 15, 20\}$ and $k' \in [5, 10]$, respectively. The hyperparameters controlling regularization strengths are $\lambda_1, \lambda_2, \lambda_3 \in [0.1, 0.3]$ with interval 0.05 for individual modality, $d_{\min} \in \{1.0, 1.5, 2.0\}$ for time series, $d_{\min} \in \{3.0, 3.5, 4.0\}$ for text. Learning rate for multi-modal encoder $\in \{0.0001, 0.0003, 0.001\}$, using Adam [84] as the optimizer. The number of case-based explanations fed to prediction LLM $\omega \in \{3, 5, 8, 10\}$. All results are reported as the average of five experimental runs.

³<https://www.kaggle.com/datasets/selfishgene/historical-hourly-weather-data>

⁴<https://www.indexmundi.com/commodities>

⁵<https://www.cdc.gov/fluview/overview/index.html>

A.3 Large Language Model

We employed the gpt-4o-2024-08-06 version for GPT-4o in OpenAI API by default. We use the parameters max_tokens=2048, top_p=1, and temperature=0.7 for content generation (self-reflection and text refinement), and 0.3 for prediction. We keep the same setting for Gemini-2.0-Flash and GPT-4o-mini due to the best empirical performance.

A.4 Environment

We conducted all the experiments on a TensorEX server with 2 Intel Xeon Gold 5218R Processor (each with 20 Core), 512GB memory, and 4 RTX A6000 GPUs (each with 48 GB memory).

B More Ablation Studies and Model Analysis

Ablations of learning objectives: We provide an ablation study on the learning objectives of TimeXL encoder, as shown in Figure 6. The results clearly show that the full objective consistently achieves the best encoder prediction performance, highlighting the necessity of regularization terms that enhance the interpretability of multi-modal prototypes. The clustering (λ_1) and evidencing (λ_2) objectives also play a crucial role in accurate prediction: the clustering term ensures distinguishable prototypes across different classes, while the evidencing term ensures accurate projection onto training data.

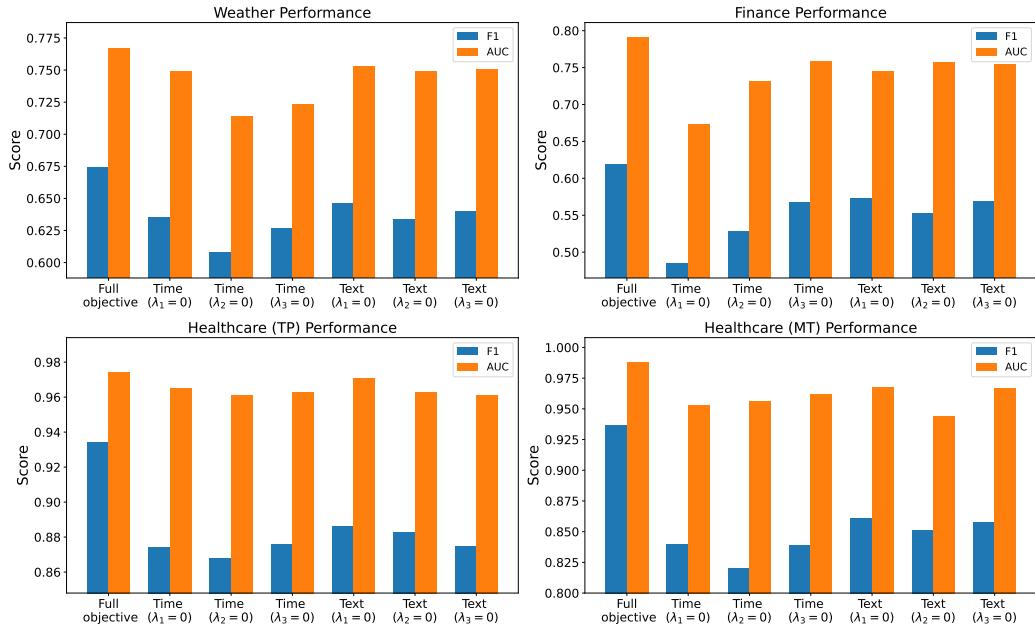


Figure 6: Ablation study of multi-modal encoder learning objectives.

Sensitivity of regularization strengths: To quantitatively evaluate the sensitivity of regularization strengths in the learning objectives, we report the performance of the multi-modal encoder on the Weather dataset (using the refined test texts), as detailed in Appendix A.2. As shown in Table 4, model performance remains relatively stable across different values. For both modalities, moderate values of λ_1 and λ_2 tend to yield better performance, suggesting the clustering and evidencing effectively guide meaningful prototype learning. A slight diversity constraint (λ_3) also helps to keep a compact interpretation prototype space.

Sensitivity of prototype lengths: We evaluate the sensitivity of the multi-modal encoder to prototype length by varying the single-kernel sizes for both time-series and text modalities on the Weather and Finance datasets (using the refined test texts), as shown in Table 5 and 6. We observe that performance is relatively stable across a range of prototype lengths on the Weather dataset, with optimal results

(larger model size & pre-training corpus). Both GPT-4o-mini and Gemini-2.0-Flash are still competitive compared with baselines listed in Table 1. It reveals the impact of the LLM capability on the effectiveness of our framework, especially in tasks demanding real-world context understanding. We also provide the plot of iterative analysis in Figure 8, where we can also observe the performance improvement over iterations for different base LLMs.

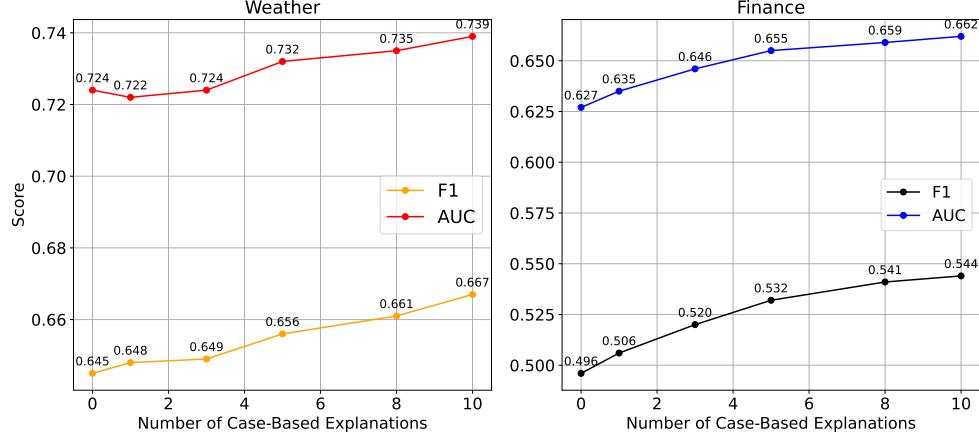


Figure 7: Effect of case-based explanations on LLM prediction performance.

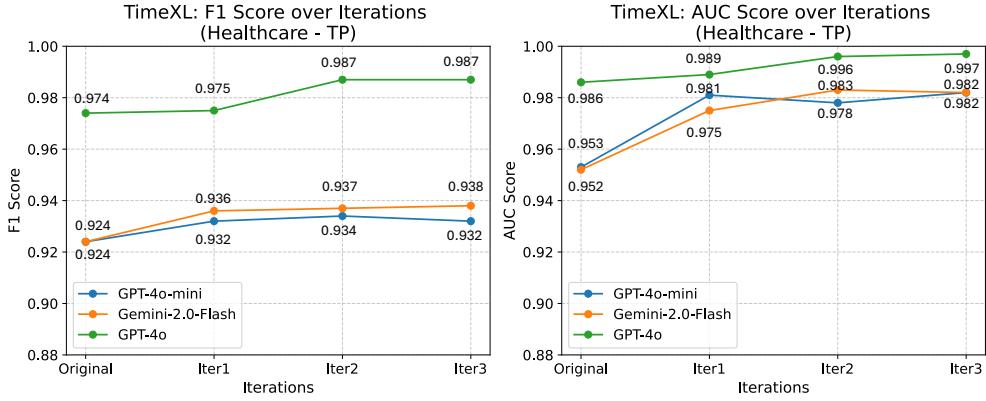


Figure 8: Iterative study of different base LLMs on the Healthcare (Test-Positive) dataset.

C Computation Cost for Iterative Optimization and Inference

We quantitatively analyze the token usage and runtime of the iterative process. Following the prompts and strategy detailed in Appendix E, each iteration over 4000 training and validation samples (average length of 400 tokens, around 300 words) requires about 9.74 M input tokens and yield 1.64 M output tokens. With GPT-4o (with input latency as $TTFT \approx 0.45$ seconds and output speed ≈ 156 tokens/second), it corresponds to around 3.9 hours serially per iteration (prediction takes approximately 0.5 hours, reflection 0.08 hours, and refinement 3.35 hours). As prediction and refinement steps can also be parallelized, we can reduce the iteration time to about 12 minutes with 20 concurrent calls. Using a more efficient LLM base model Gemini-Flash-2.0 with $TTFT \approx 0.3$ seconds and ≈ 236 output tokens/second, the iteration time can be further reduced to 8 minutes.

In the testing stage, we apply the fixed reflection selected during validation to refine the test texts. Based on our prompt lengths, the refinement and prediction steps take approximately 3.46 seconds per sample using GPT-4o, and 2.29 seconds per sample for Gemini-2.0-Flash, which is reasonable for deployment in real-world applications.

In addition to efficient LLM backbones and parallelized querying, we discuss further strategies to lower computational cost. As shown in Figure 5, 1-2 iterations often yield clear performance gains. In practice, the iteration loop can be adaptively controlled by monitoring validation improvements and stopping early once gains fall below a predefined threshold. To further reduce the cost, we can also do selective refinement by skipping training examples that were already predicted accurately in previous iterations. These strategies help make our method more computationally tractable.

D Explainable Multi-modal Prototypes and Case Study

D.1 Multi-modal Prototypes for All Datasets

We present the learned multi-modal prototypes across datasets including Weather (Figure 9), Finance (Figure 10), Healthcare (Test-Positive) (Figure 11), and Healthcare (Mortality) (Figure 12). It is noticeable that the prototypes from both modalities align well with real-world ground truth scenarios, ensuring faithful explanations and enhancing LLM predictions.

D.2 Case-based Reasoning Example on the Finance Dataset

We provide another case-based reasoning example to demonstrate the effectiveness of TimeXL in explanatory analysis, as shown in Figure 13. In this example, the original text is incorrectly predicted as neutral instead of a decreasing trend of raw material prices. We have a few key observations based on the results. First, the refinement LLM filters the original text to emphasize economic and market conditions more indicative of a declining trend, based on the reflections from training examples. The refined text preserves discussions on port inventories and steel margins while placing more emphasis on subdued demand, thin profit margins, and bearish market sentiment as key indicator of prediction. Accordingly, the case-based explanations from the original text focus more on inventory management and short-term stable patterns, while those in the refined text highlight demand contraction, production constraints, and macroeconomic uncertainty, which is more consistent with a decreasing trend. Furthermore, the reasoning on time series provides a complementary view for predicting raw material price trends. The time series explanations identify declining price movements across multiple indices. In general, the multi-modal explanations based on matched prototypes from TimeXL demonstrate its effectiveness in capturing relevant raw material market condition for both predictive and explanatory analysis.

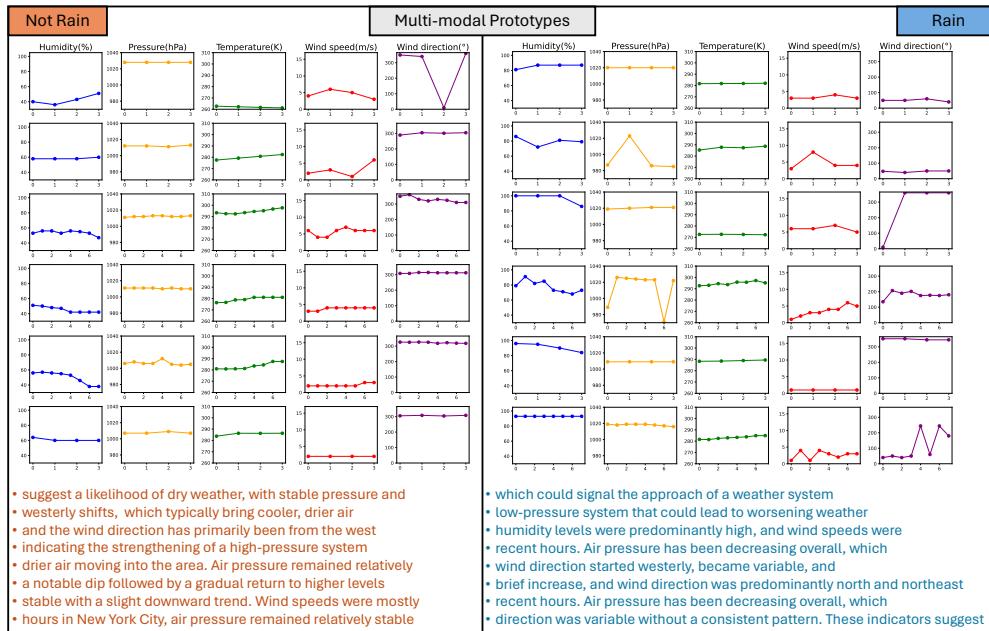


Figure 9: More multi-modal prototypes learned from the Weather dataset. Each row in the figure represents a time series prototype.

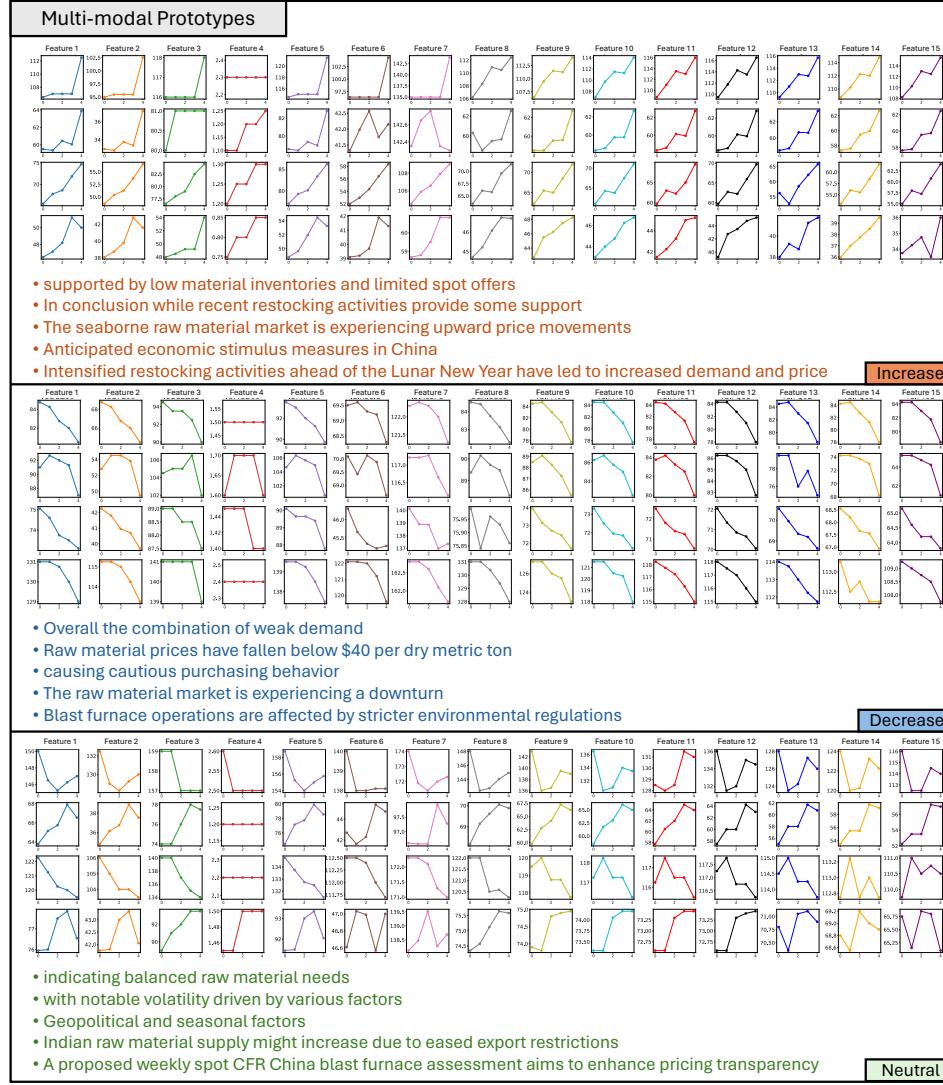


Figure 10: Key multi-modal prototypes learned from the Finance dataset. Each row in the figure represents a time series prototype.

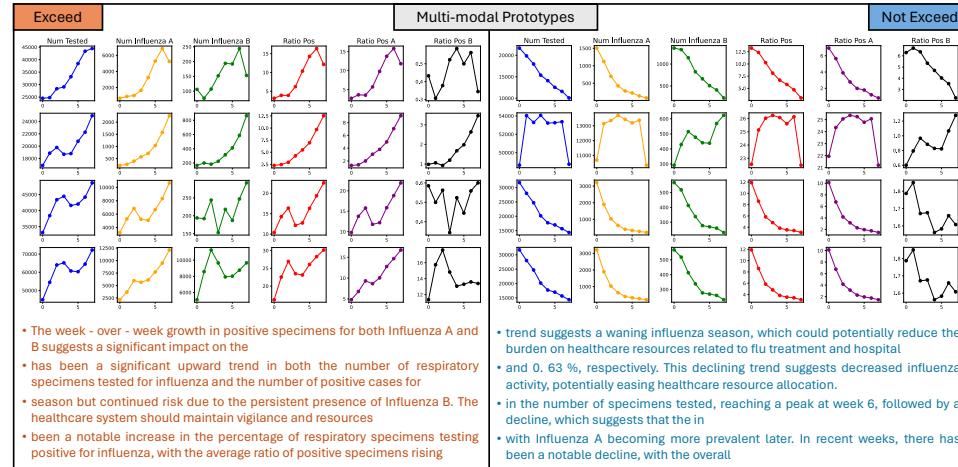


Figure 11: Key multi-modal prototypes learned from the Healthcare (Test-Positive) dataset. Each row in the figure represents a time series prototype.

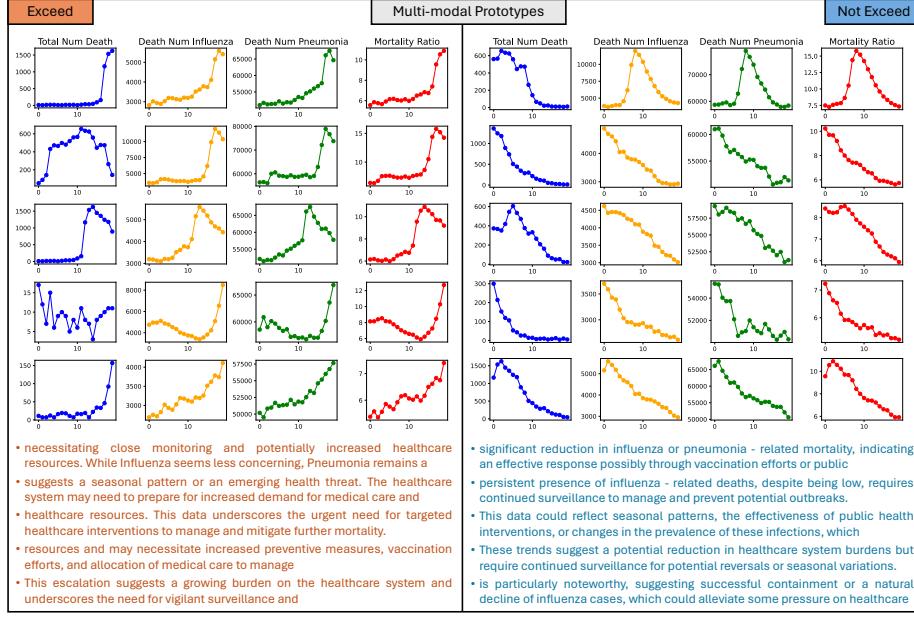


Figure 12: Key multi-modal prototypes learned from the Healthcare (Mortality) dataset. Each row in the figure represents a time series prototype.

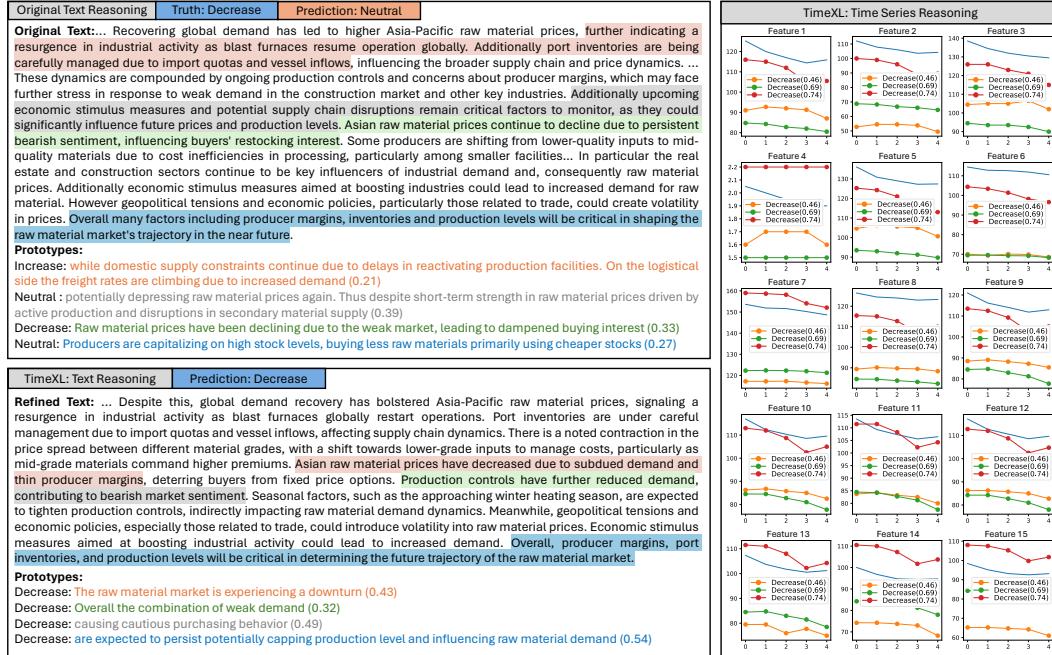


Figure 13: Multi-modal case-based reasoning example on the Finance dataset.

E Designed Prompts for Experiments

In this section, we provide our prompts for prediction LLM in Figure 14 (and a text-only variant for comparison, in Figure 15), reflection LLM in Figures 16, 17 18, as well as refinement LLM in Figure 19. Note that we adopt a generate-update-summarize strategy to effectively capture the reflective thoughts from training samples with class imbalances, which is more structured and scalable. We make the whole training texts into batches. First, the reflection LLM generates the initial reflection (Figure 16) by extracting key insights from class-specific summaries, highlighting text

patterns that contribute to correct and incorrect predictions. Next, it updates the reflection (Figure 17) by incorporating new training data, ensuring incremental and context-aware refinements. Finally, it summarizes multiple reflections from each class (in Figure 18) into a comprehensive guideline for downstream text refinement. This strategy consolidates knowledge from correct predictions while learning from incorrect ones, akin to the training process of deep models.

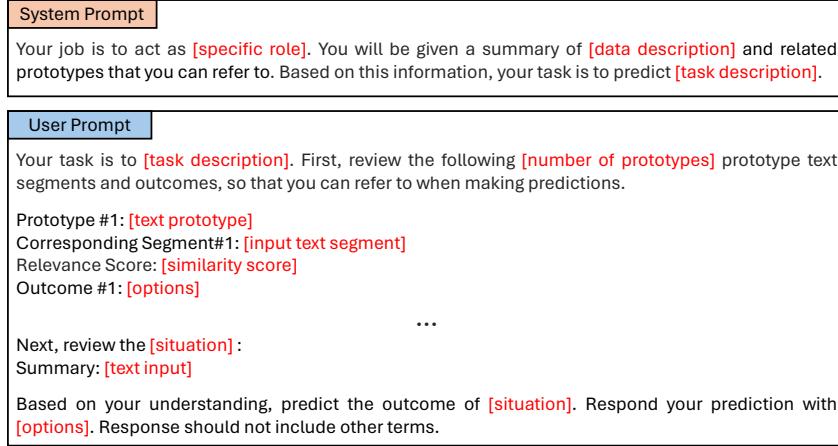


Figure 14: Prompt for prediction LLM

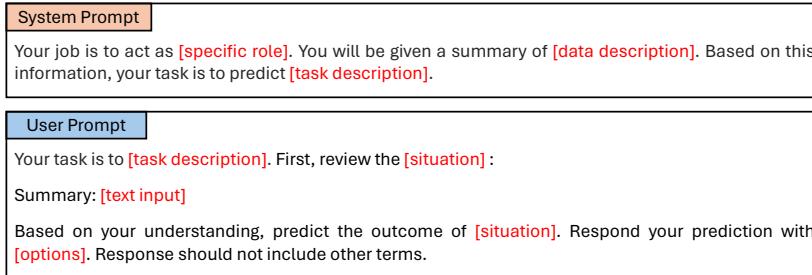


Figure 15: Prompt for prediction with text only

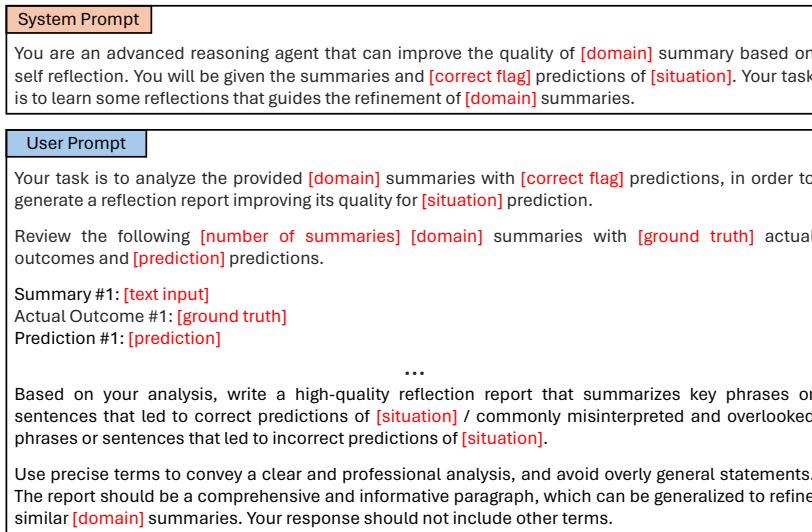


Figure 16: Prompt for reflection LLM: reflection generation

| |
|--|
| System Prompt |
| <p>You are an advanced reasoning agent that can improve the quality of [domain] summary based on self reflection. You will receive a reflection report up to this point. You will also be given the summaries and [correct flag] predictions of [situation]. Your task is to learn some reflections and update the current report that guides the refinement of [domain] summaries.</p> |
| User Prompt |
| <p>Your task is to analyze the provided [domain] summaries with [correct flag] predictions, in order to update a reflection report improving its quality for [situation] prediction.</p> <p>First, review the following reflection report up to this point: [current reflection report]</p> <p>Next, review the following [number of summaries] [domain] summaries with [ground truth] actual outcomes and [prediction] predictions.</p> <p>Summary #1: [text input] Actual Outcome #1: [ground truth] Prediction #1: [prediction]</p> <p>...</p> <p>Based on your analysis, write a high-quality reflection report that summarizes key phrases or sentences that led to correct predictions of [situation] / commonly misinterpreted and overlooked phrases or sentences that led to incorrect predictions of [situation].</p> <p>Use precise terms to convey a clear and professional analysis, and avoid overly general statements. The report should contain incremental and context-aware updates, and can be generalized to refine similar [domain] summaries. Your response should not include other terms.</p> |

Figure 17: Prompt for reflection LLM: reflection update

| |
|---|
| System Prompt |
| <p>You are an advanced summarization agent that can generate high-quality summarization. You will be given previously generated reflections for text refinement, from the correct and incorrect predictions of [domain] texts. Your current task is to summarize these long reflections to better guide financial text refinement.</p> |
| User Prompt |
| <p>Your task is to summarize the long reflections derived from previous predictions of [domain] contents. The goal is to generate a high-quality report aimed at improving the [domain] text quality for better predictive accuracy.</p> <p>First, review the reflections from all combinations of possible predictions and actual outcomes: [reflection reports]</p> <p>Based on your analysis, summarize the reflections of different scenarios and write a comprehensive report that provides guidelines to select the most important content in new [domain] texts where the actual outcome is unknown. Your response should keep the enough details, yet effective, to improve the text quality for downstream prediction. Your response should not include other terms.</p> |

Figure 18: Prompt for reflection LLM: reflection summarization

| |
|---|
| System Prompt |
| <p>You are an advanced refinement agent designed to enhance the quality of [domain] summary. You will be provided with reflective thoughts analyzed from other summaries, and a summary that requires refinement. Your task is to generate a refined [domain] summary, by examining how reflective thoughts applied to the current summary.</p> |
| User Prompt |
| <p>Your task is to generate a refined weather summary from the current summary to improve its predictions of [situation]. First, review the following reflections that provide guidelines for refinement: [final reflection report]</p> <p>Next, review the current [domain] summary that describes [situation]:</p> <p>Summary #1: [text input]</p> <p>Based on your understanding, generate a new weather summary by selecting relevant content in the current summary, which provides insights crucial for understanding [situation]. Response should not include other terms.</p> |

Figure 19: Prompt for refinement LLM

F Reflection Reports for Text Refinement

In this section, we provide the reflection reports after the first iteration, to demonstrate the reflective thoughts by accessing real-world contexts.

Reflection Summary – Weather (New York)

Comprehensive Report on Refining Weather Predictions

The review of weather summaries across various predictive scenarios reveals critical insights that can enhance the accuracy and reliability of weather forecasts. By integrating these refined elements, future weather texts can better anticipate actual outcomes. Below are the summarized guidelines for improving weather prediction texts:

1.Key Indicators for Dry Conditions ("Not Rained" Predictions and Outcomes):

1. **Stable and Slightly Increasing Air Pressure:** Consistent readings with minor upward trends suggest high-pressure systems, indicative of dry weather.
2. **Gentle to Moderate Wind Speeds:** Observations of stable wind speeds without gusts support non-precipitative forecasts.
3. **Variable Wind Directions (Northwesterly/Westerly):** Shifts from southwesterly to northwesterly/westerly directions bring cooler, drier air, reducing rain likelihood.
4. **Decreasing Daytime Humidity:** High early humidity followed by daytime decreases correlates with dry conditions.
5. **Typical Diurnal Temperature Patterns:** Normal temperature variations further support dry forecasts.
6. **Absence of Significant Weather Systems:** Lack of major air pressure or wind pattern changes reinforces stable, dry conditions.

2.Indicators for Rainy Conditions ("Rained" Predictions and Outcomes):

1. **Increasing Humidity Levels:** Significant humidity increases are strong rain indicators.
2. **Decreasing Air Pressure:** A downward pressure trend signals potential rain due to incoming low-pressure systems.
3. **Wind Conditions:**
 1. **Slight Wind Speed Increase:** Often precedes rain, particularly if observed later.
 2. **Easterly/Southeasterly Wind Directions:** Bring moisture-laden air, favoring rain.

3.Common Misinterpretations ("Rained" Predictions with "Not Rained" Outcomes):

1. **Pressure Stability Misinterpretation:** Minor pressure fluctuations without other indicators are weak rain predictors.
2. **Humidity and Fog Confusion:** Differentiate between humidity peaks indicating fog versus those suggesting rain.
3. **Overemphasis on Wind Changes:** Focus on sustained wind patterns rather than minor variations.
4. **Misjudged Temperature Fluctuations:** Evaluate temperature changes alongside corroborative indicators like pressure drops.
5. **Dynamic Weather System Misinterpretation:** Recognize stable atmospheric conditions to avoid false rain predictions.

4.Overlooked Conditions ("Not Rained" Predictions with "Rained" Outcomes):

1. **Humidity Peaks and Variability:** High peaks or fluctuations often precede rain; they shouldn't be underestimated.
2. **Wind Direction Shifts:** Changes to southerly or easterly directions can signal moisture influx.
3. **Temperature and Humidity Interactions:** Combined changes may offset stable pressure, indicating rain potential.
4. **Stable Pressure with Other Variables:** Consider pressure stability with high humidity and wind changes collectively.

By focusing on these detailed indicators and refining the interpretation of weather elements, future weather texts can improve prediction quality, providing more accurate and reliable forecasts.

Figure 20: Reflection summary for text refinement on the Weather dataset.

Reflection Summary – Finance (Raw Material)

Comprehensive Guidelines for Enhancing Financial Text Quality in Predictive Analysis

To improve the accuracy of financial predictions, it's crucial to refine the selection and emphasis of content in financial summaries. Based on reflections from past predictions, the following guidelines highlight the essential elements to consider for optimizing text quality.

1.Key Indicators for Price Trends:

1. For predicting increases, focus on indicators such as robust demand from major consumers, strong profit margins in key manufacturing sectors, and preferences for high-grade raw materials. Also, consider tight supply chains, low inventories, and economic stimulus measures.
2. For decreases, emphasize rising inventory levels, weak demand, and operational adjustments. Bearish market sentiment and geopolitical tensions should be noted as well.
3. For neutral outcomes, identify market equilibrium indicators like balanced supply-demand dynamics, stable production operations, and moderate buyer behavior.

2.Contextual Influences:

1. Consider the impact of seasonal trends, restocking activities, and economic stimuli on demand fluctuations. Recognize how these factors may cause temporary market shifts rather than long-term trends.
2. Analyze global trade and supply chain disruptions to assess their potential to cause short-term volatility or stability rather than sustained changes.

3.Regulatory and Policy Factors:

1. Understand the implications of environmental regulations and geopolitical policies on supply and demand. These elements can significantly alter market dynamics and should be carefully integrated into analyses.

4.Market Sentiment and Related Markets:

1. Assess market sentiment and futures movements, recognizing that positive futures alignments often indicate underlying demand. Consider interconnected markets, such as related raw materials sectors, for their influence on overall demand.

5.Strategic and Operational Adjustments:

1. Pay attention to how producers and buyers adjust operations and procurement strategies in response to market conditions. These adjustments can affect demand patterns and price stability.

6.Avoiding Common Misinterpretations:

1. Avoid overemphasizing short-term trends as indicators of sustained market shifts. Be wary of speculative influences and temporary policy announcements that may not have immediate or lasting effects.
2. Recognize high inventory levels as potential strategic buffers and consider their role in stabilizing prices during off-peak periods.

By integrating these guidelines, financial summaries can provide a more comprehensive and balanced analysis, enhancing the predictive accuracy of unknown market outcomes. The focus should remain on understanding and articulating the complexities of market dynamics to better align predictions with actual trends.

Figure 21: Reflection summary for text refinement on the Finance dataset.

Reflection Summary – Healthcare (Test-Positive)

Comprehensive Report on Improving Healthcare Text Quality for Predictive Accuracy

Introduction

This report synthesizes reflections from past analyses of healthcare summaries concerning influenza positivity rates. The goal is to enhance the quality of these texts to improve predictive accuracy for future assessments where actual outcomes are unknown.

Key Indicators for Accurate Predictions

1.Declining Trends for "Not Exceed" Predictions

1. **Emphasis on Decline:** Correct predictions for outcomes that did not exceed the average often highlighted a clear decline in positivity rates. Terms such as "notable decline," "steady decrease," and "significant reduction" are crucial.
2. **Comparative Analysis:** A pronounced reduction in Influenza A and B, particularly Influenza A, supports accurate assessments.
3. **Testing Volume and Ratios:** A decrease in testing volume, coupled with positivity rates remaining below the average, strengthens predictions.
4. **Low Positivity Percentages:** Explicit references to recent positivity rates consistently falling below the long-term average are vital.

2.Indicators for "Exceed" Predictions

1. **Marked Increases:** Substantial increases in positivity rates, highlighted by specific numerical comparisons, indicate an "exceed" outcome.
2. **Influenza Strain Dominance:** Details on the predominance of Influenza A or B, with their impact on overall positivity rates, are significant.
3. **Rising Testing Volumes:** Increased specimen testing, peaking with positivity rates, suggests heightened incidence and surveillance.
4. **Healthcare System Impact:** References to increased hospitalizations and medical service demand provide context and validate exceeding predictions.
5. **Temporal Patterns:** Tracking weekly peaks aids in understanding and predicting heightened influenza activity.

Common Misinterpretations

- Relative vs. Absolute Increases:** Misinterpretations often arise from conflating relative increases with exceeding the average positivity rate. It is essential to compare these increases explicitly with historical averages.
- Contextual Misalignment:** Phrases highlighting rises in positivity rates without contextual alignment to the average threshold can mislead predictions.

Guidelines for Text Refinement

- Integrate Clear Comparisons:** Ensure summaries incorporate direct comparisons to historical averages to avoid misjudgments.
- Maintain Trend Clarity:** Differentiate clearly between increasing trends and those that actually surpass the average positivity rate.
- Include Contextual Analysis:** Provide context for increases in positivity rates by detailing the impact on healthcare resources and system strain.
- Highlight Temporal Data:** Incorporate temporal patterns and weekly peaks to aid in nuanced predictive analysis.

By adhering to these guidelines, healthcare texts can be refined to enhance predictive accuracy, ultimately aiding in proactive healthcare planning and resource allocation.

Figure 22: Reflection summary for text refinement on the Healthcare (Test-Positive) dataset.

| Reflection Summary – Healthcare (Mortality) | |
|---|--|
| Reflection Summary Report for Healthcare Text Refinement: | |
| 1.Correct 'Not Exceed' Predictions: | |
| 1. Key indicators of accurate predictions include phrases such as "general downward trend," "consistent decline," and "significant decrease" in mortality ratios, particularly from Influenza or Pneumonia. | |
| 2. Terms like "below the average ratio" and "remained under the historical benchmark" effectively signal outcomes that do not surpass average mortality ratios. | |
| 3. References to "improvement in management" and "effective public health interventions" highlight successful control measures that contribute to lower mortality figures. | |
| 4. Acknowledging a "persistent burden" of Pneumonia amidst declines provides a nuanced understanding of stability in disease impact. | |
| 2.Correct 'Exceed' Predictions: | |
| 1. Accurate forecasts are marked by phrases such as "concerning upward trend," "significant increase," and "notable spike," indicating rising mortality ratios. | |
| 2. Quantitative expressions like "dramatic rise" in deaths and comparisons exceeding specific averages effectively contextualize high mortality trends. | |
| 3. Descriptions of peaks, potential outbreaks, and healthcare strain underscore the urgency and resource implications, enhancing predictive accuracy. | |
| 3.Incorrect 'Exceed' Predictions with 'Not Exceed' Outcomes: | |
| 1. Misinterpretations arise from emphasizing upward trends without aligning them with final ratios below the 7.84% threshold. | |
| 2. Overemphasis on isolated peaks without considering their unsustained nature leads to overestimations. | |
| 3. For improved accuracy, texts should differentiate between temporary increases and sustained trends, and focus on cumulative ratios relative to average benchmarks. | |
| 4.Incorrect 'Not Exceed' Predictions with 'Exceed' Outcomes: | |
| 1. Underestimations often occur by not fully considering the impact of recent historical peaks or seasonal outbreaks, even when declines are present. | |
| 2. Incorrect predictions result from neglecting cumulative effects of influenza and pneumonia, especially when one disease overshadows the other. | |
| 3. Future accuracy can be enhanced by emphasizing recent peaks, seasonal patterns, and comprehensive disease impacts. | |
| Guidelines for Future Healthcare Texts: | |
| •Prioritize clear identification of trends and ratios, ensuring they are contextualized against historical averages and thresholds. | |
| •Distinguish between short-term fluctuations and sustained trends, emphasizing cumulative impacts over isolated data points. | |
| •Integrate insights from public health interventions and management effectiveness to provide a holistic view of disease control. | |
| •Consider seasonal patterns and recent historical data to anticipate potential outbreaks or shifts in disease prevalence. | |
| •Ensure quantitative comparisons are explicit and supported by comprehensive data analysis to avoid misinterpretations. | |
| This structured approach will aid in crafting accurate healthcare summaries, enhancing predictive accuracy and improving the quality of healthcare texts for future analyses. | |

Figure 23: Reflection summary for text refinement on the Healthcare (Mortality) dataset.

G TimeXL for Regression-based Prediction: A Finance Demonstration

In this section, we provide a demonstration of using TimeXL for regression tasks. Two minor modifications adapt TimeXL for continuous value prediction. First, we add a regression branch in the encoder design, as shown in Figure 24. On the top of time series prototype layers, we reversely ensemble each time series segment representation as a weighted sum of time series prototypes, and add a regression head (a fully connected layer with non-negative weights) for prediction. Accordingly, we add another regression loss term (*i.e.*, MSE and MAE) to the learning objectives. Second, we adjust the prompt for prediction LLM by adding time series inputs and requesting continuous forecast values, as shown in Figure 25. As such, TimeXL is equipped with regression capability.

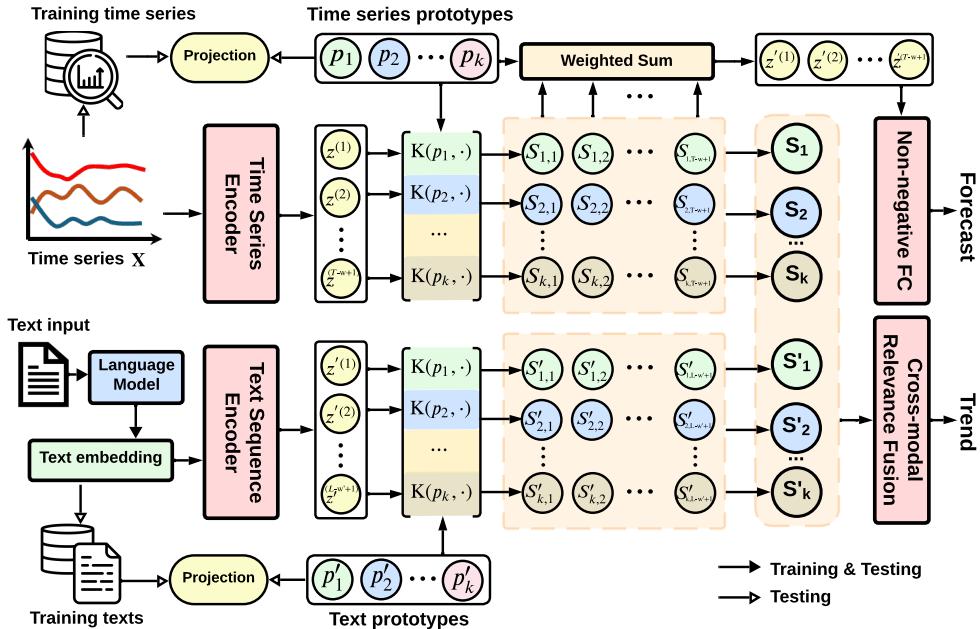


Figure 24: Multi-modal prototype-based encoder design in TimeXL for regression tasks.

| |
|--|
| System Prompt |
| Your job is to act as [specific role]. You will be given a summary of [data description] and related prototypes that you can refer to. Based on this information, your task is to predict [task description]. |
| User Prompt |
| Your task is to [task description]. First, review the following [number of prototypes] prototype text segments and outcomes, so that you can refer to when making predictions. |
| <p>Prototype #1: [text prototype] Corresponding Segment#1: [input text segment] Relevance Score: [similarity score] Outcome #1: [options]</p> <p style="text-align: center;">...</p> <p>Next, review the [situation] : Summary: [text input]</p> <p>Finally, review the [domain] record of [situation] : [time series values]</p> <p>Based on your understanding, predict the outcome of [situation], followed by the value of [domain] record. Respond your prediction with [options] and [numerical value]. Response should not include other terms.</p> |

Figure 25: Prompt for prediction LLM in regression tasks.

Table 8: Performance evaluation for regression tasks.

| Model | RMSE | MAE | MAPE(%) |
|-----------------------------|--------------|--------------|--------------|
| DLinear [75] | 7.871 | 6.400 | 4.727 |
| Autoformer [76] | 7.215 | 5.680 | 4.263 |
| Crossformer [77] | 7.205 | 5.313 | 3.808 |
| TimesNet [78] | 6.978 | 4.928 | 3.512 |
| iTransformer [79] | 5.877 | 4.023 | 2.863 |
| TSMixer [80] | 7.447 | 5.509 | 3.911 |
| TimeMixer [81] | 6.651 | 4.703 | 3.349 |
| FreTS [82] | 7.098 | 4.886 | 3.460 |
| PatchTST [5] | 5.676 | 4.042 | 2.853 |
| LLMTTime [83] | 11.545 | 5.300 | 3.774 |
| PromptCast [63] | 4.728 | 3.227 | 2.306 |
| OFA [53] | 6.906 | 4.862 | 3.463 |
| Time-LLM [56] | 6.396 | 4.534 | 3.238 |
| TimeCMA [57] | 7.187 | 5.083 | 3.620 |
| FSCA [58] | 5.511 | 3.873 | 2.720 |
| MM-iTransformer [35] | 5.454 | 3.789 | 2.687 |
| MM-PatchTST [35] | 5.117 | 3.493 | 2.491 |
| TimeCAP [65] | 4.456 | 3.088 | 2.196 |
| TimeXL | 4.161 | 2.844 | 2.035 |

We conduct the experiments on the same Finance dataset, where the target value is the raw material stock price. The performance of all baselines and TimeXL is presented in Table 8. Consistent with the classification setting, TimeXL achieves the best results, and the multi-modal variants of state-of-the-art baselines (MM-iTransformer and MM-PatchTST) benefit from incorporating additional text data. This observation is further supported by the ablation study in Table 9. We also note that the text modality offers complementary information for LLMs, enhancing numerical price predictions. Furthermore, the identified prototypes provide contextual guidance, leading to additional performance gains. Overall, TimeXL outperforms all variants, underscoring the effectiveness of mutually augmented prediction. We also provide an iteration analysis to show the effectiveness of reflective and refinement process, as shown in Table 10. The prediction performance quickly improves and stabilizes over iterations, which underscores the alternation steps between predictions and reflective refinements.

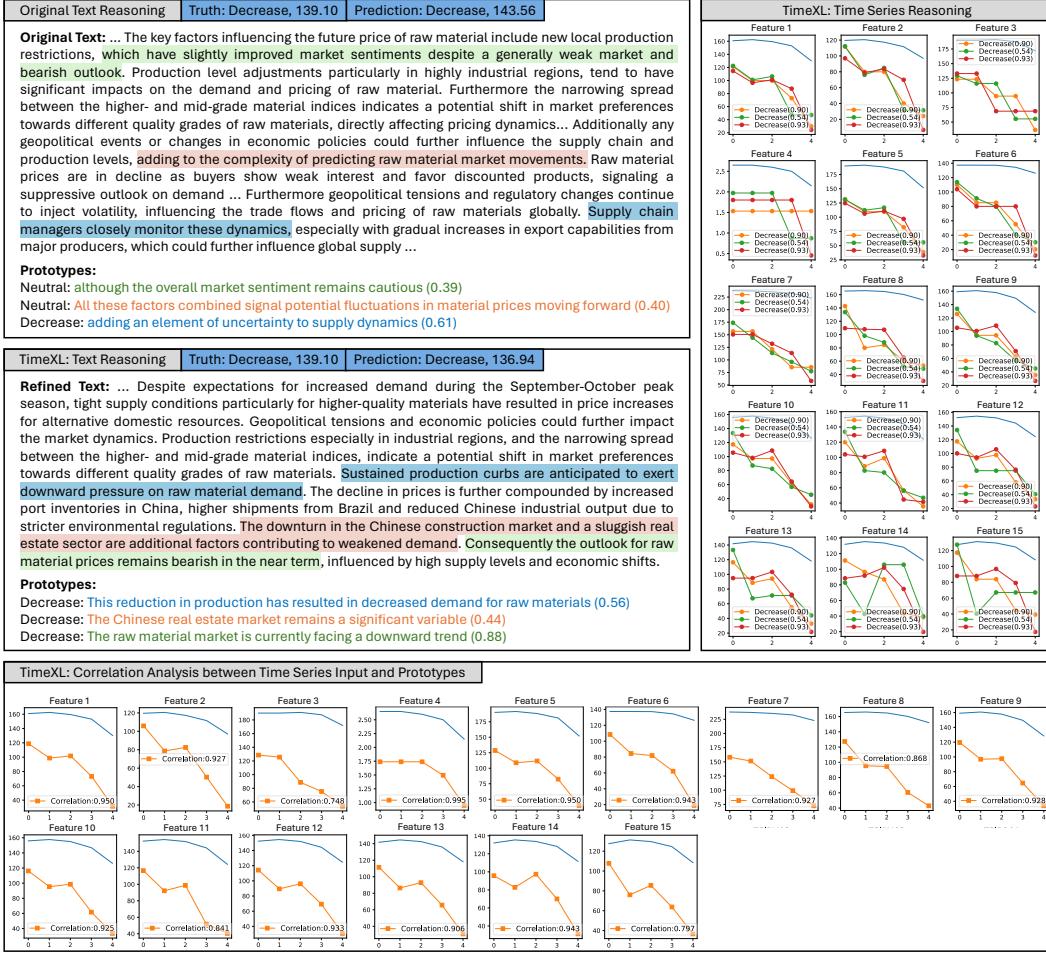


Figure 26: Multi-modal case-based reasoning example on the Finance dataset for regression tasks.

Table 9: Ablation studies of TimeXL for regression tasks.

| Variants | RMSE | MAE | MAPE(%) |
|---|--------------|--------------|--------------|
| Time Series Prototype-based Encoder | 4.287 | 3.001 | 2.140 |
| Multi-modal Prototype-based Encoder | 4.198 | 2.891 | 2.064 |
| Prediction LLM - Time Series + Text | 4.600 | 3.121 | 2.226 |
| Prediction LLM - Time Series + Text + Prototype | 4.352 | 3.003 | 2.165 |
| TimeXL | 4.161 | 2.844 | 2.035 |

Table 10: Iterative analysis of TimeXL for regression tasks.

| Iteration | RMSE | MAE | MAPE(%) |
|-----------|-------|-------|---------|
| Original | 4.344 | 2.951 | 2.103 |
| 1 | 4.224 | 2.883 | 2.069 |
| 2 | 4.161 | 2.844 | 2.035 |
| 3 | 4.174 | 2.849 | 2.036 |

Finally, we provide a case study on the test set of the Finance dataset that demonstrates the effectiveness of TimeXL’s reasoning process for regression tasks, as shown in Figure 26. We present

the top three most similar prototypes allocated for both time series and text modalities, where we have two key observations. 1) Even if the trend predictions align correctly in both original and refined texts, the prototype allocation varies regarding market sentiments. In the original texts, prototypes from the neutral class are also allocated, highlighting a mixed sentiment and potential fluctuations, consequently deteriorating numerical prediction accuracy from both the encoder and the LLM. In the refined texts, prototypes instead capture the underlying bearish market sentiment for raw materials, emphasizing decreased demand and weakened market conditions, effectively aligning bearish market sentiment with numerical predictions. 2) The allocated top similar time series prototypes also exhibit highly similar trends compared to the original input series, offering robust reference points for numerical prediction. To quantitatively assess it, we aggregate the retrieved prototypes using their similarity scores and compute the correlation between the input time series and the aggregated prototype series. The observed strong correlations underscores TimeXL’s efficacy in capturing intrinsic temporal patterns and providing predictive insights into future numerical stock values based on similar historical patterns.

H More Detailed Discussions of Related Work

Due to space limitations, we provide additional details on related methods here, expanding on the overview in the main text.

H.1 Multi-modal Time Series Analysis.

In recent years, multi-modal time series analysis has gained significant traction in diverse domains such as finance, healthcare and environmental sciences [28, 11, 13]. Multiple approaches have been proposed to model interactions across different modalities, such as multi-modal fusion and cross-modal alignment and so on. Multi-modal fusion captures complementary information to enhance time series modeling, which are commonly implemented via addition [35] and concatenation [29] at the multiple levels. For example, Liu *et al.* [35] fuse predictions from both time series and text modalities via learnable linear weighted addition, further enhancing the predictive performance of state-of-the-art forecasters. In parallel, cross-modal alignment aims to capture the association between time series and other modalities for downstream tasks, where common techniques include attention mechanism [12, 30, 86, 31], gating functions [32], and contrastive learning [33, 34]. For instance, Bamford *et al.* [14] align time series and text in a shared latent space using deep encoders, enabling retrieval of specific time series sequences based on textual queries. In addition, Zheng *et al.* [33] perform causal structure learning in multi-modal time series by separating modality-invariant and modality-specific components through contrastive learning, enabling root cause analysis. Besides method developments, recent research efforts have also advanced the field through comprehensive benchmarks and studies [35, 36, 37], highlighting the incorporation of new modalities to boost task performance. Meanwhile, others explore representing time series as other modalities (*e.g.*, image, text, table, graph) [38, 39, 40, 41, 42] or generating synthetic modalities [32, 43] to facilitate downstream tasks. Nevertheless, these studies tend to focus primarily on improving numerical accuracy. The deeper and tractable rationale behind *why* or *how* the textual or other contextual signals influence time series outcomes remains underexplored.

H.2 Time Series Explanation

Recent studies have explored diverse paradigms for time series interpretability. Gradient-based and perturbation-based explanations leverage saliency methods to highlight important features at different time steps [44, 45], where some methods explicitly incorporate temporal structures into models and objectives [46, 47]. Meanwhile, surrogate approaches also offer global or local explanations by evaluating the importance of time series features with respect to prediction. These methods include the generalization of Shapley value to time series [48], the formulation of model behavior consistency via self-supervised learning [49], and the usage of information-theoretic guided objectives for coherent explanations [50]. In contrast to saliency or surrogate-based explanations, we adopt a *case-based reasoning* paradigm [25, 26, 27], which end-to-end generates predictions and built-in explanations from learned prototypes. Our work extends this approach to multi-modal time series by producing human-readable reasoning artifacts for both the temporal and contextual modalities.

H.3 LLMs for Time Series Analysis

The rapid development of LLMs [15, 16] has begun to inspire new directions in time series research [51, 52]. Many existing methods fine-tune pre-trained LLMs on time series tasks and achieving state-of-the-art results in forecasting, classification, and beyond [53, 54, 55]. The textual data (*e.g.*, domain instructions, metadata, and dataset summaries) are often encoded as prefix embeddings to enrich time series representations [56, 57, 58, 59]. These techniques also contribute to the emergence of time series foundation models [55, 60, 61, 62]. An alternative line of research leverages the zero-shot or few-shot capabilities of LLMs by directly prompting pre-trained language models with text-converted time series [63] or context-laden prompts representing domain knowledge [24], often yielding surprisingly strong performance in real-world scenarios. Furthermore, LLMs can act as knowledge inference modules, synthesizing high-level patterns or explanations that augment standard time series pipelines [64, 65, 66]. For instance, Wang *et al.* [66] construct a large news database for method development and perform text-to-text prediction by fine-tuning an LLM with dynamically selected news, which uses reflective selection logic to incorporate matched social events and thus augment numerical prediction. Building on this trend, recent works have explored time series reasoning with LLMs by framing tasks as natural-language problems [68], including time series understanding & question answering (*e.g.*, trend and seasonal pattern recognition, similarity comparison, causality analysis, *etc.*) [67, 69, 70, 71, 72], as well as language-guided inference for standard zero-shot time series tasks [67, 69]. Compared to existing LLM-based methods, our approach provides a unique synergy between an explainable model and LLMs through iterative, grounded, and reflective interaction. It explicitly provides case-based explanations from both modalities, highlighting key time series patterns and refined textual segments that contribute to the prediction.